

# CS2040S

## Data Structures and Algorithms

### Shortest Paths!

---


(Try writing a program to solve this!)

#### Puzzle of the week:

- 5 x 5 grid
- Choose a starting square
- Move: 3 cells vertically or horizontally OR
- Move: 2 cells diagonally.
- Cannot visit same cell twice.
- Cannot exit grid


To win: visit all cells.

Example:



\*\* What's the *worst* you can do?

# SHORTEST PATHS



# What is a directed graph?

---

Graph consists of two types of elements:

Nodes (or vertices)

- At least one.


Edges (or arcs)

- Each edge connects two nodes in the graph
- Each edge is unique.
- Each edge is **directed**.

# What is a directed graph?

---

Graph  $G = \langle V, E \rangle$

- $V$  is a set of nodes
    - At least one:  $|V| > 0$ .
  - $E$  is a set of edges:
    - $E \subseteq \{ (v,w) : (v \in V), (w \in V) \}$
    - $e = (v,w)$
    - For all  $e_1, e_2 \in E : e_1 \neq e_2$
- Order matters! 

# What is a directed graph?


---

In-degree: number of incoming edges

Out-degree: number of outgoing edges


Out-degree: 3

In-degree: 2  
Out-degree: 1



Is it a directed graph?


- ✓ 1. Yes
- 2. No.



Is it a directed graph?


1. Yes

✓ 2. No.



Is it a directed graph?

- ✓ 1. Yes
- 2. No.





# Representing a (Directed) Graph

---

## Adjacency List:

- Array of nodes
- Each node maintains a list of neighbors
- Space:  $O(V + E)$



## Adjacency Matrix:

- Matrix  $A[v,w]$  represents edge  $(v,w)$
- Space:  $O(V^2)$

# Adjacency List

Graph consists of:

- Nodes: stored in an array
- Edges: linked list per node





# Adjacency List

---

Directed Graph consists of:

- Nodes: stored in an array
- **Outgoing** Edges: linked list per node




# Adjacency List in Java

---

```
class NeighborList extends ArrayList<Integer> {  
}
```

```
class Node {  
    int key;  
    NeighborList nbrs;  
}
```

```
class Graph {  
    Node[] nodeList;  
  
}
```



# Representing a (Directed) Graph

---

## Adjacency List:

- Array of nodes
- Each node maintains a list of neighbors
- Space:  $O(V + E)$

## Adjacency Matrix:

- Matrix  $A[v,w]$  represents edge  $(v,w)$
- Space:  $O(V^2)$


# Adjacency Matrix

---

Graph consists of:

- Nodes
- Edges = pairs of nodes

	a	b	c	d	e	f
a	0	0	0	0	1	1
b	0	0	1	1	1	0
c	0	1	0	0	0	0
d	0	1	0	0	0	0
e	1	1	0	0	0	0
f	1	0	0	0	0	0




# Adjacency Matrix

---

Directed Graph consists of:

- Nodes
- Edges = pairs of nodes

	a	b	c	d	e	f
a	0	0	0	0	1	1
b	0	0	1	1	0	0
c	0	1	0	0	0	0
d	0	1	0	0	0	0
e	0	1	0	0	0	0
f	0	0	0	0	0	0



# Adjacency Matrix

---

Graph represented as:

$$A[v][w] = 1 \text{ iff } (v,w) \in E$$

	a	b	c	d	e	f
a	0	0	0	0	<b>1</b>	<b>1</b>
b	0	0	<b>1</b>	<b>1</b>	<b>0</b>	0
c	0	<b>1</b>	0	0	0	0
d	0	<b>1</b>	0	0	0	0
e	<b>0</b>	<b>1</b>	0	0	0	0
f	<b>0</b>	0	0	0	0	0



# Searching a (Directed) Graph

---

## Breadth-First Search:

- Search level-by-level
- Follow outgoing edges
- Ignore incoming edges

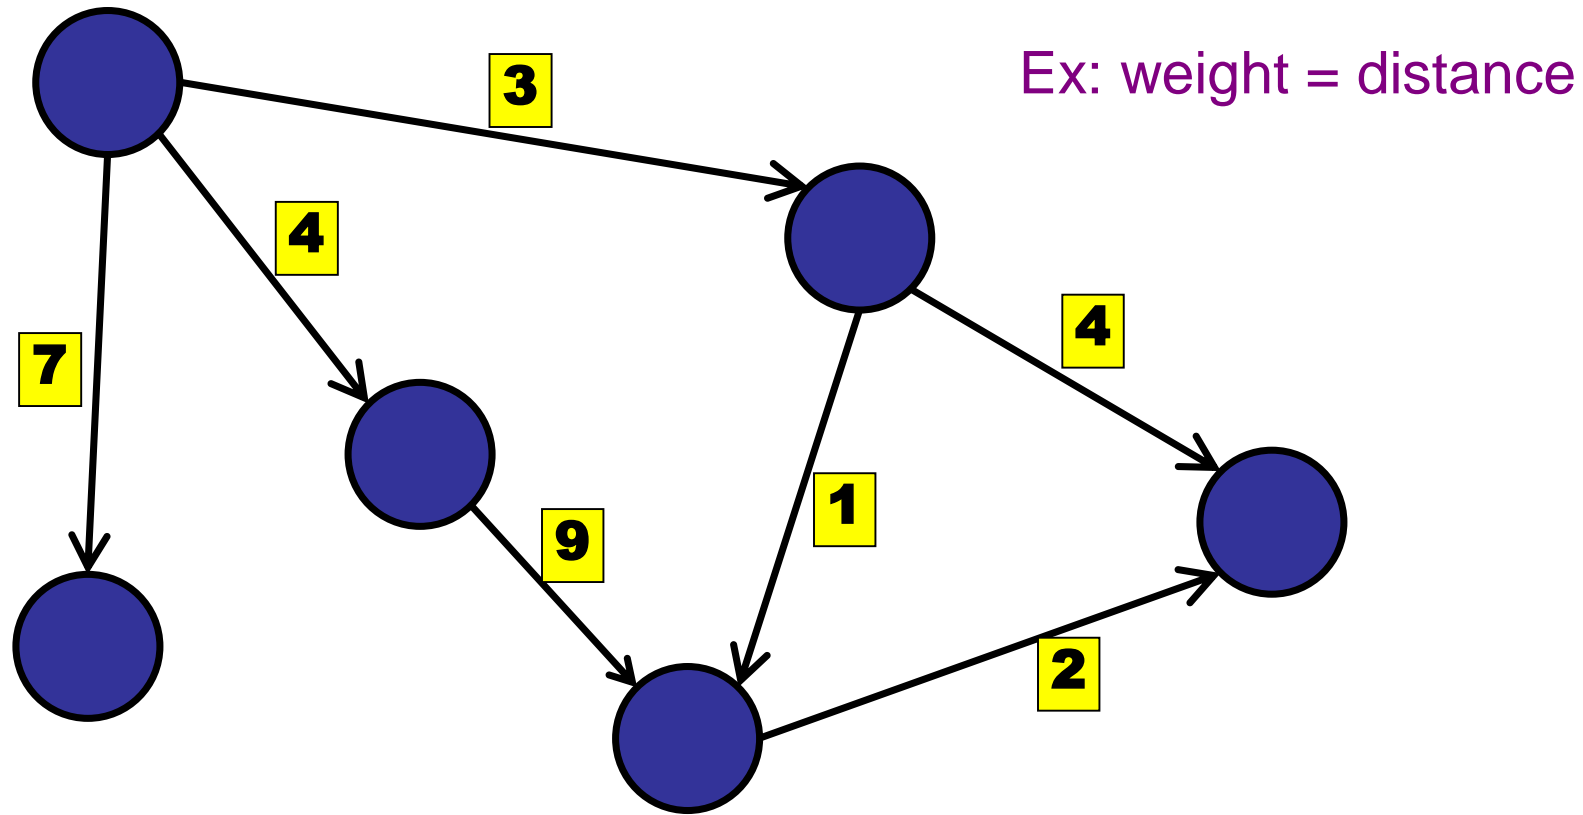
## Depth-First Search:

- Search recursively
- Follow outgoing edges
- Backtrack (through incoming edges)

# Weighted Graphs

---

**Edge weights:**  $w(e) : E \rightarrow \mathbb{R}$




Adjacency list: stores weights with edge in NbrList

# Shortest Paths


---

Distance from source?



What is the distance from S to D?


- 1. 2
- 2. 4
- ✓ 3. 6
- 4. 7
- 5. 9
- 6. Infinite



# Shortest Paths

---

Distance from source?




# Shortest Paths

---

## Questions:


- How far is it from S to D?
- What is the shortest path from S to D?
- Find the shortest path from S to every node.
- Find the shortest path between every pair of nodes.



# Shortest Paths

---


Common mistake: "Why can't I use BFS?"



# Shortest Paths

---


Common mistake: "Why can't I use BFS?"





# Shortest Paths

Common mistake: "Why can't I use BFS?"




BFS finds minimum number of **HOPS** not minimum **DISTANCE**.

# Shortest Paths

---

Notation:  $\delta(u,v)$  = distance from u to v




# Shortest Paths

---

Key idea: triangle inequality

$$\delta(S, C) \leq \delta(S, A) + \delta(A, C)$$



# Shortest Paths


---

Maintain estimate for each distance:

```
int[] dist = new int[V.length];
```

```
Arrays.fill(dist, INFTY);
```

```
dist[start] = 0;
```




# Shortest Paths

---

Maintain estimate for each distance:

- Reduce estimate
- Invariant: estimate  $\geq$  distance




# Shortest Paths

---

Maintain estimate for each distance:


$\text{relax}(S, A)$



# Shortest Paths

---

```
relax(int u, int v){  
    if (dist[v] > dist[u] + weight(u,v))  
        dist[v] = dist[u] + weight(u,v);  
}
```




# Shortest Paths

---

Maintain estimate for each distance:

$\text{relax}(S, A)$





# Shortest Paths

Maintain estimate for each distance:


$\text{relax}(A, C)$

Triangle Inequality:

$$\delta(S, C) \leq \delta(S, A) + \delta(A, C)$$



Can safely reduce estimate at C to 7.




# Shortest Paths

---

Maintain estimate for each distance:

$\text{relax}(A, C)$




# Shortest Paths

---

Maintain estimate for each distance:

$\text{relax}(A, B)$




# Shortest Paths

---

Maintain estimate for each distance:

$\text{relax}(S, B)$




# Shortest Paths

---

Maintain estimate for each distance:


`relax(B, C)`



# Shortest Paths

---

```
for (Edge e : graph)
    relax(e)
```

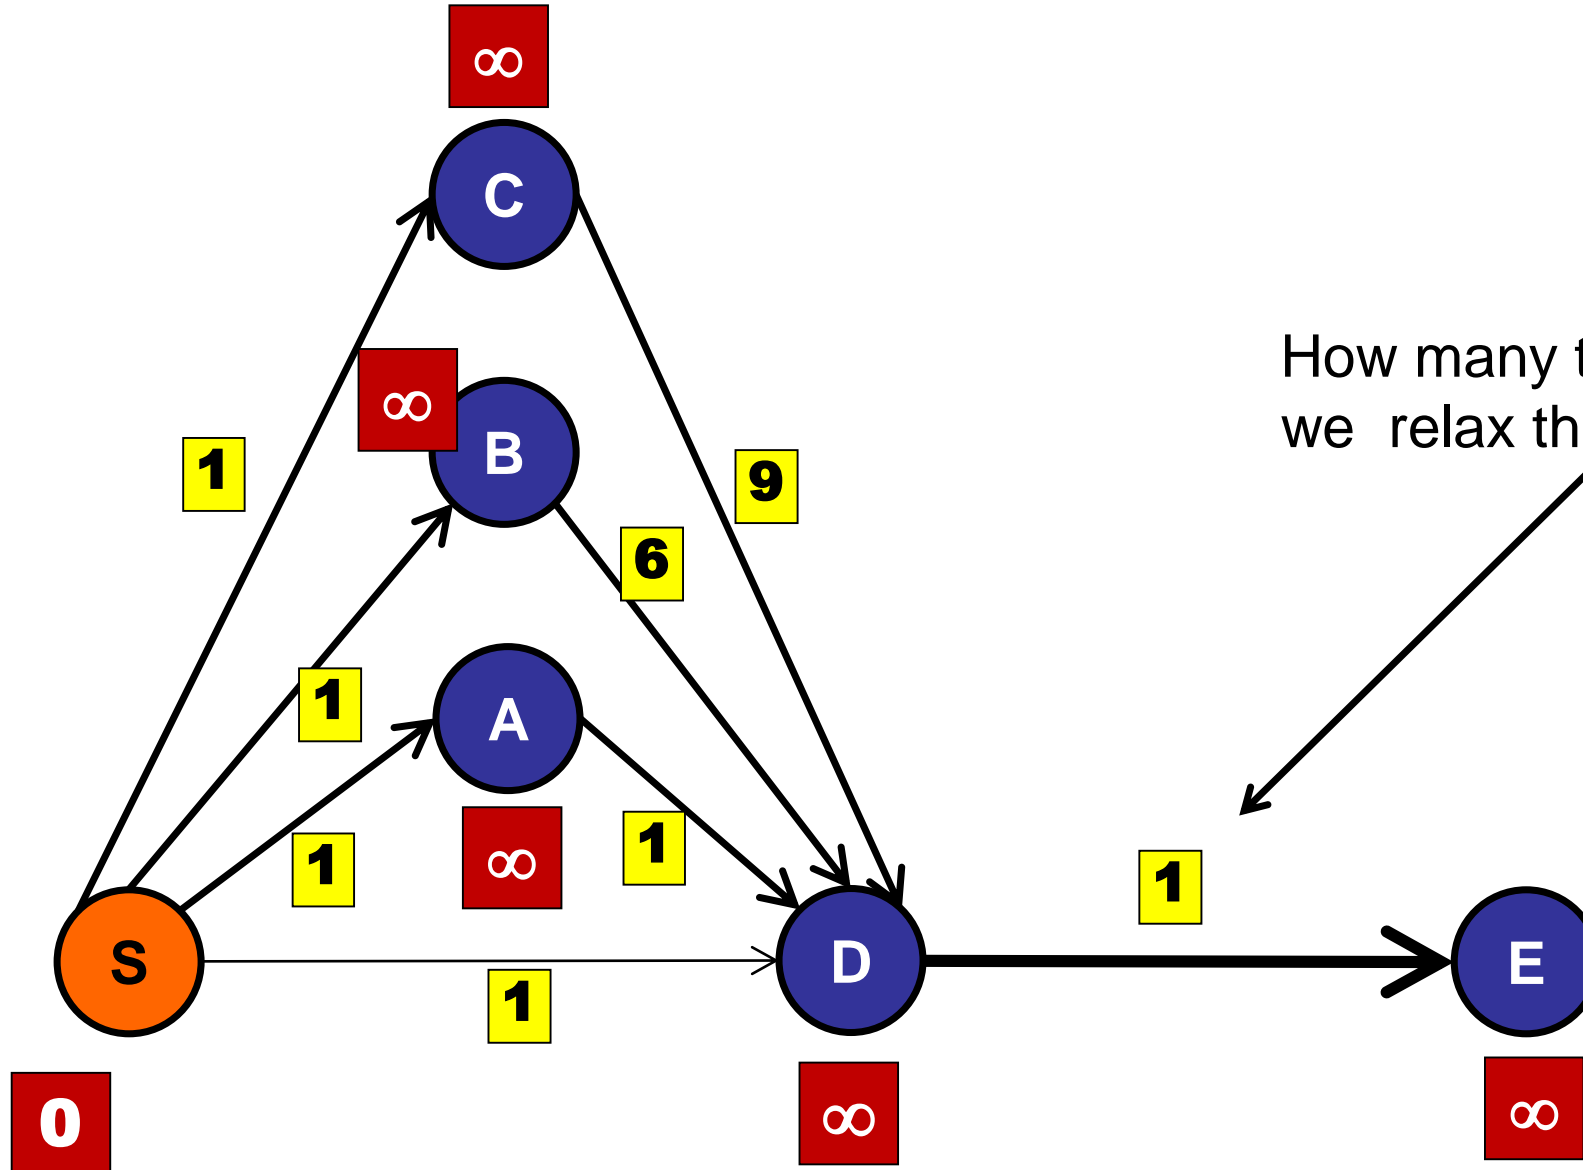


Does this algorithm work:

**for every edge  $e$ : relax( $e$ )**

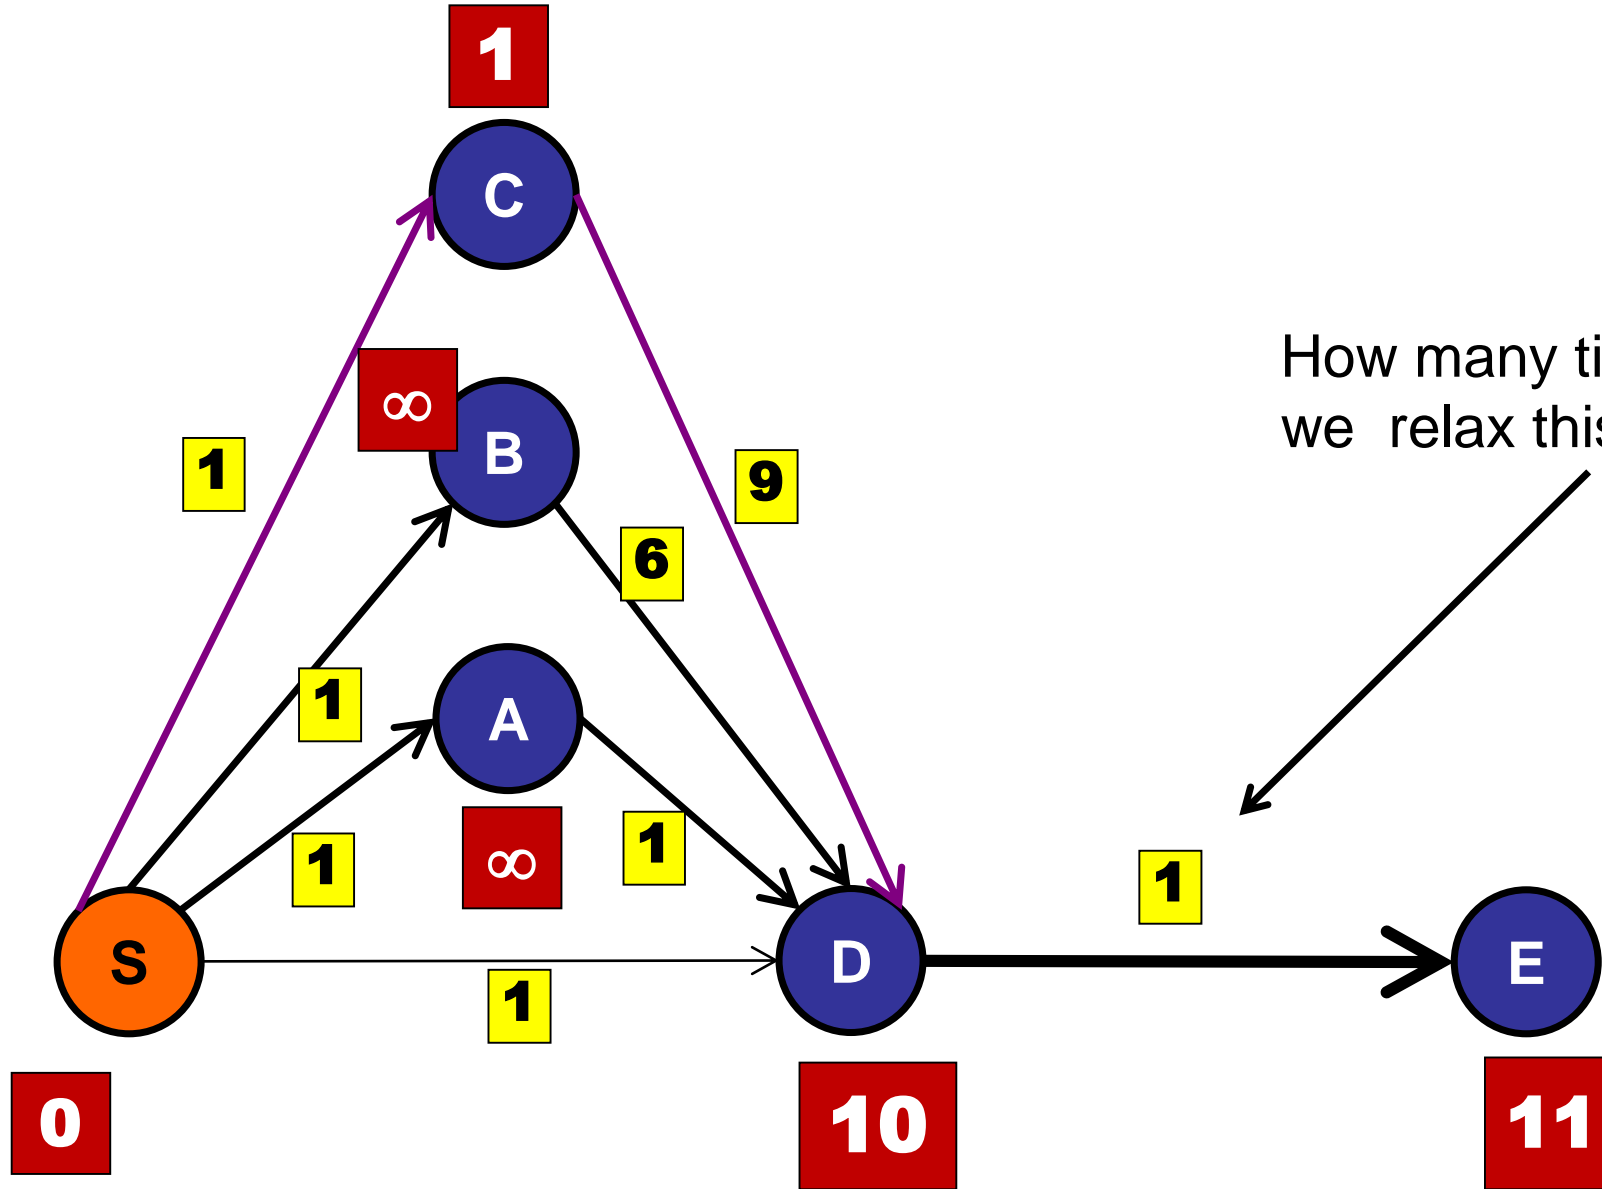
1. Yes
2. Sometimes
3. No

# Shortest Paths



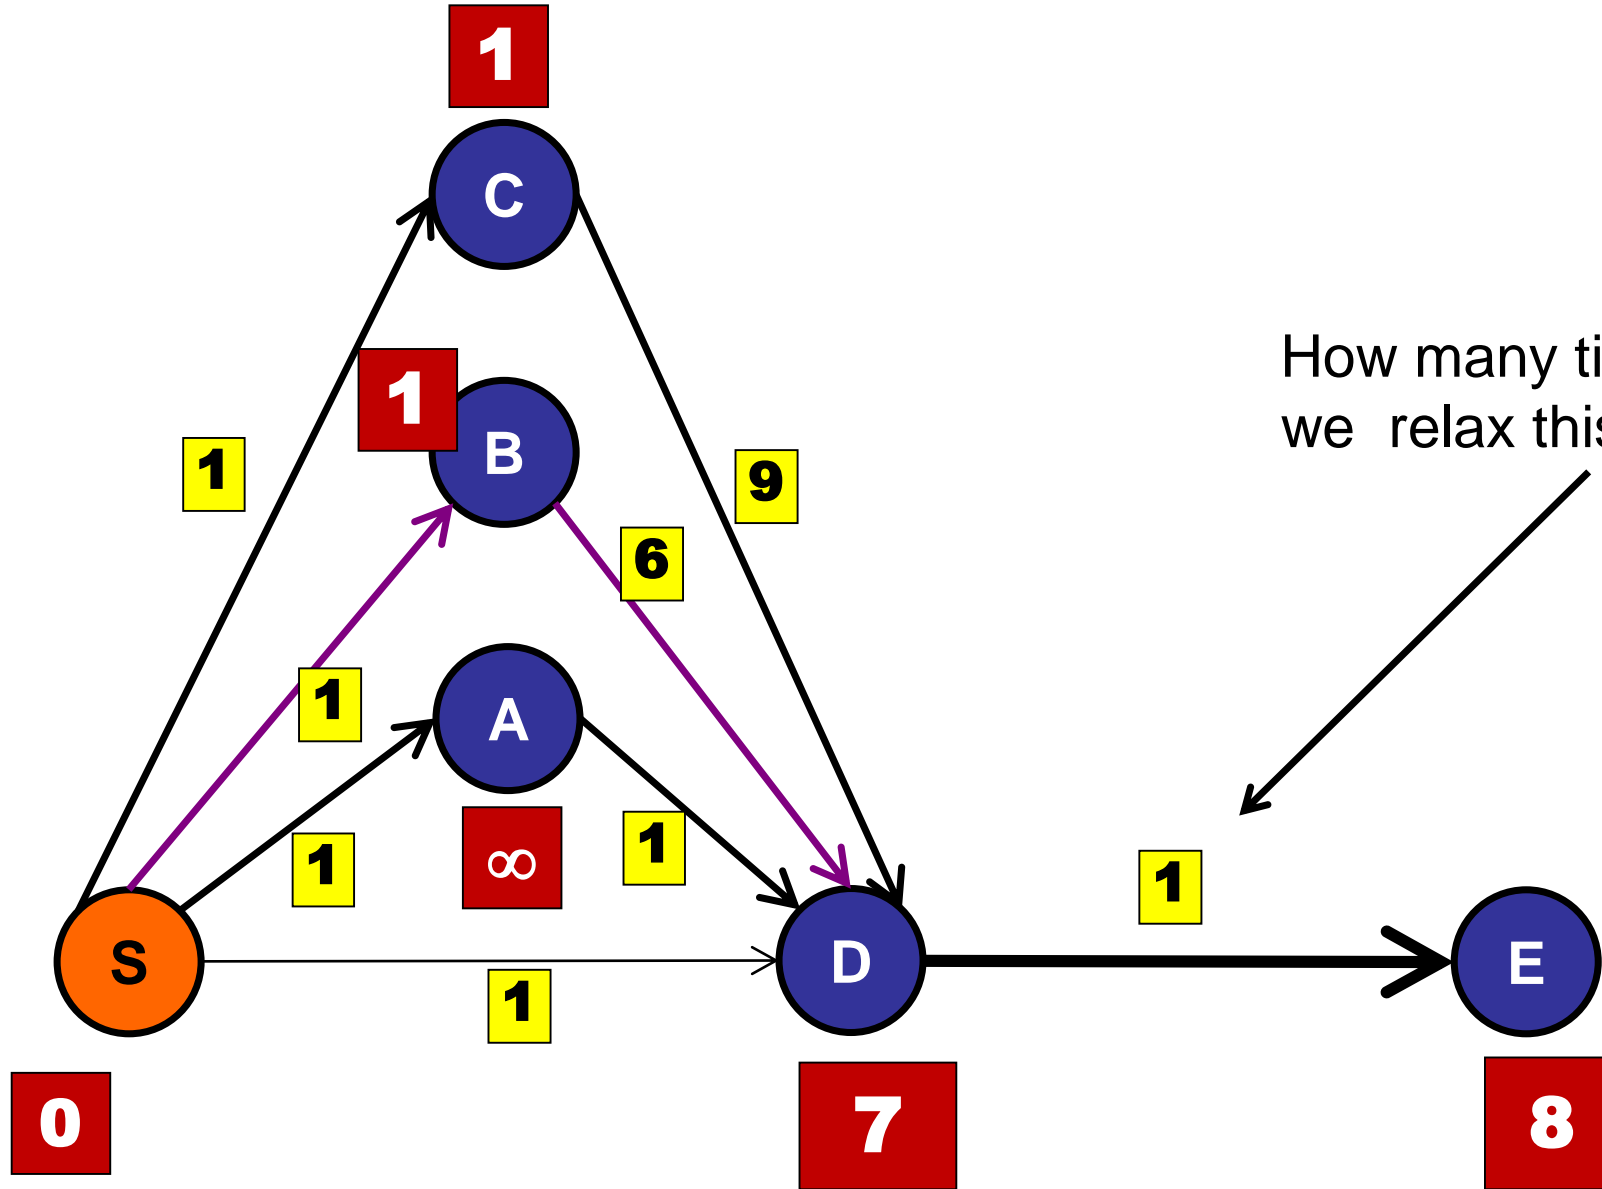


# Shortest Paths

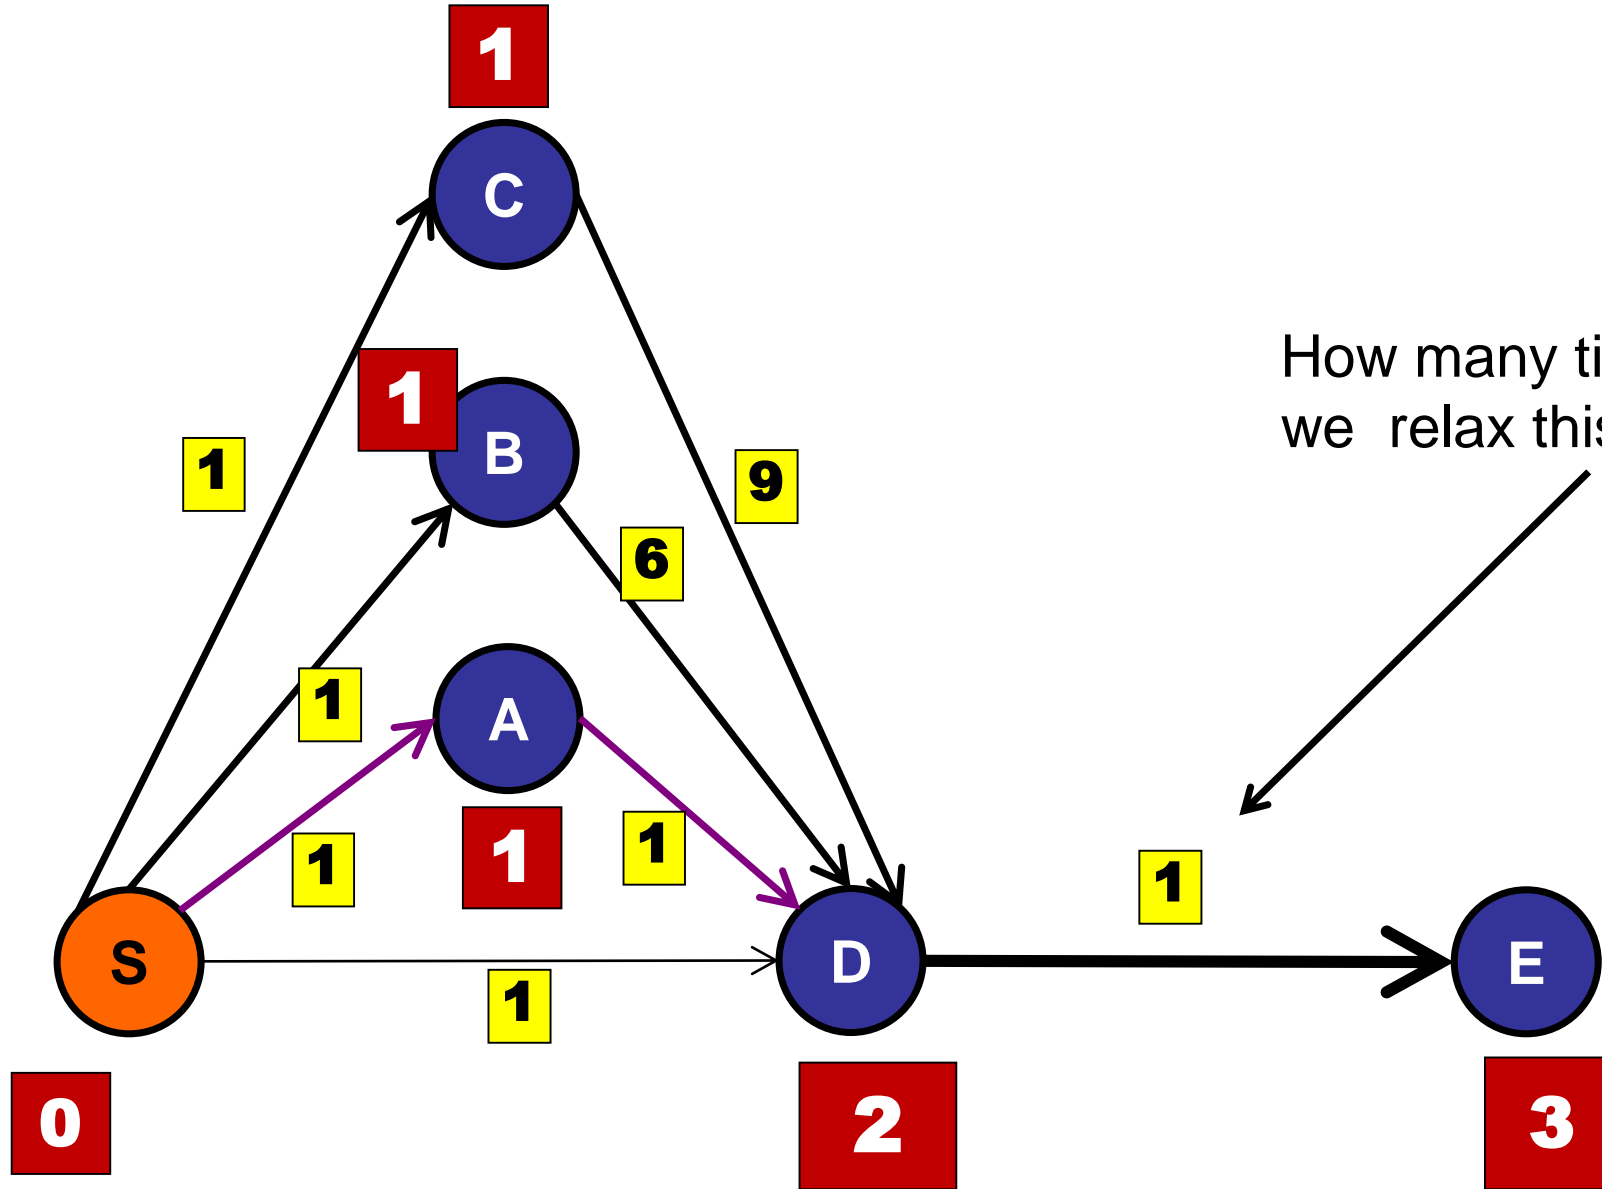


How many times might we relax this edge?

# Shortest Paths

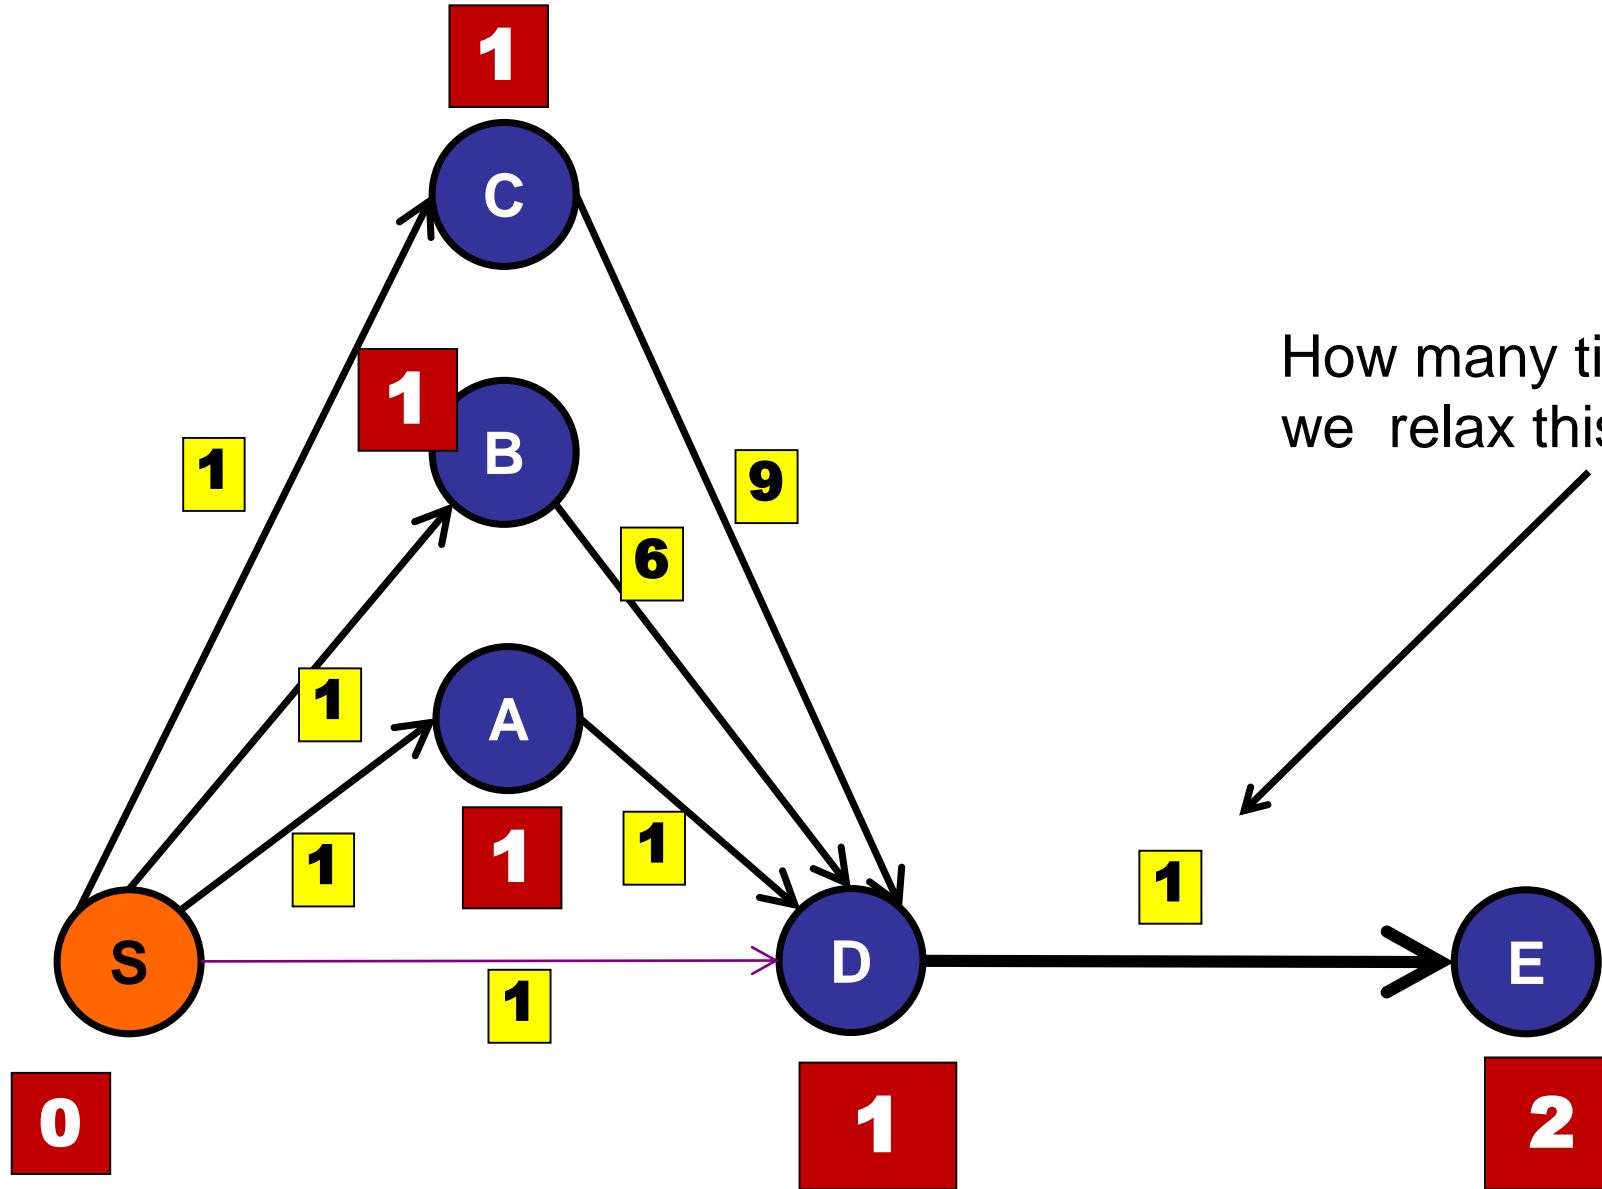


# Shortest Paths



How many times might we relax this edge?

# Shortest Paths



How many times might we relax this edge?

# Bellman-Ford


```
n = V.length;  
for (i=0; i<n; i++)  
    for (Edge e : graph)  
        relax(e)
```




Richard Bellman



Lester R. Ford, Jr




## When can you terminate early?

1. When a relax operation has no effect.
2. When two consecutive relax operations have no effect.
-  3. When an entire sequence of  $|E|$  relax operations have no effect.
4. Never. Only after  $|V|$  complete iterations.

# Bellman-Ford

---

```
n = V.length;  
  
for (i=0; i<n; i++)  
    for (Edge e : graph)  
        relax(e)
```



What is the running time of Bellman-Ford?


1.  $O(V)$
2.  $O(E)$
3.  $O(V+E)$
4.  $O(E \log V)$
- ✓ 5.  $O(EV)$



# Bellman-Ford

---

```
n = V.length;  
  
for (i=0; i<n; i++)  
    for (Edge e : graph)  
        relax(e)
```



# Bellman-Ford

---

## Properties:

- $O(EV)$  running time (in the worst-case)
- Can stop after one entire iteration with no changes to the estimates.

## Invariant:

- Let  $T$  be a *shortest path tree* of graph  $G$  rooted at source  $s$ .
- After iteration  $j$ , if node  $u$  is  $j$  hops from  $s$  on tree  $T$ , then  $est[u] = distance(s, u)$ .

# Bellman-Ford


---

Why does this work?

# Bellman-Ford

---


Why does this work?



# Bellman-Ford


---

Why does this work?



Look at minimum weight path from S to D.  
(Path is simple: no loops.)

# Why does this work?




After 1 iteration, 1 hop estimate is correct.

# Bellman-Ford

---

Why does this work?




After 1 iteration, 1 hop estimate is correct.

# Bellman-Ford

---

Why does this work?




After 1 iteration, 1 hop estimate is correct.



# Bellman-Ford

---

Why does this work?




After 2 iterations, 2 hop estimate is correct.

# Bellman-Ford

---

Why does this work?




After 3 iterations, 3 hop estimate is correct.

# Bellman-Ford

---

Why does this work?




After 4 iterations, D estimate is correct.

# Bellman-Ford

---


What if edges have negative weight?



# Bellman-Ford

---

What if edges have negative weight?




No problem!

# Bellman-Ford

---


What if edges have negative weight?



# Bellman-Ford

---

What if edges have negative weight?




$d(S,C)$  is infinitely negative!

# Negative weight cycles

---

How to detect negative weight cycles?






# Negative weight cycles

---

How to detect negative weight cycles?




Run Bellman-Ford for  $|V|+1$  iterations.

If an estimate changes in the last iteration... then negative weight cycle.

# Bellman-Ford

---


Special case: all edges have the same weight



# Bellman-Ford

---

Special case: all edges have the same weight.



Use regular Breadth-First Search.

# Bellman-Ford Summary

---

## Basic idea:

- Repeat  $|V|$  times: relax every edge
- Stop when “converges”.
- $O(VE)$  time.

## Special issues:

- If negative weight-cycle: impossible.
- Use Bellman-Ford to detect negative weight cycle.
- If all weights are the same, use BFS.

# Faster algorithms?

---

Key idea:

Relax the edges in the “right” order.

Only relax each edge once:


- $O(E)$  cost (for relaxation step).

# Relax edges in the right order

---

If there are no negative weight cycles, is there *always* a "right" order to relax the edges?

If so, prove it.



If not, give a counter-example.

**\*\* a "right" order is one where each edge is relaxed only once.**

# Faster algorithms?

---

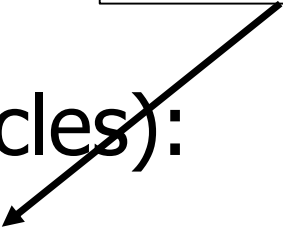
Key idea:

Relax the edges in the “right” order.

Only relax each edge once:

- $O(E)$  cost (for relaxation step).

not a useful  
algorithm!



A right order always exists (if no neg. wt. cycles):

- Find shortest path tree.
- Relax tree edges in breadth-first order.
- Relax non-tree edges in any order.

# Faster algorithms?

---

Key idea:

Relax the edges in the “right” order.


Only relax each edge once:

- $O(E)$  cost (for relaxation step).

Necessary assumption:

All edges weights  $\geq 0$ .

Extending a path does not make it shorter!






# Dijkstra's Algorithm (First Try)


---

Relax shortest edge first




# Dijkstra's Algorithm (First Try)

Relax shortest edge first




# Dijkstra's Algorithm (First Try)

Relax shortest edge first




# Dijkstra's Algorithm (First Try)

Relax shortest edge first



# Dijkstra's Algorithm (First Try)


Relax shortest edge first



# Dijkstra's Algorithm (Failed Try)

Oops....

Only relax an edge once its estimate is correct (and will never change again)!



# Dijkstra's Algorithm


---

## Basic idea:

- Maintain distance estimate for every node.
- Begin with empty shortest-path-tree.
- Repeat:
  - Consider **node** with minimum estimate.
  - (We will show that this node has a good estimate.)
  - Add node to shortest-path-tree.
  - Relax all outgoing edges.

# Shortest Paths

---








# Dijkstra's Algorithm

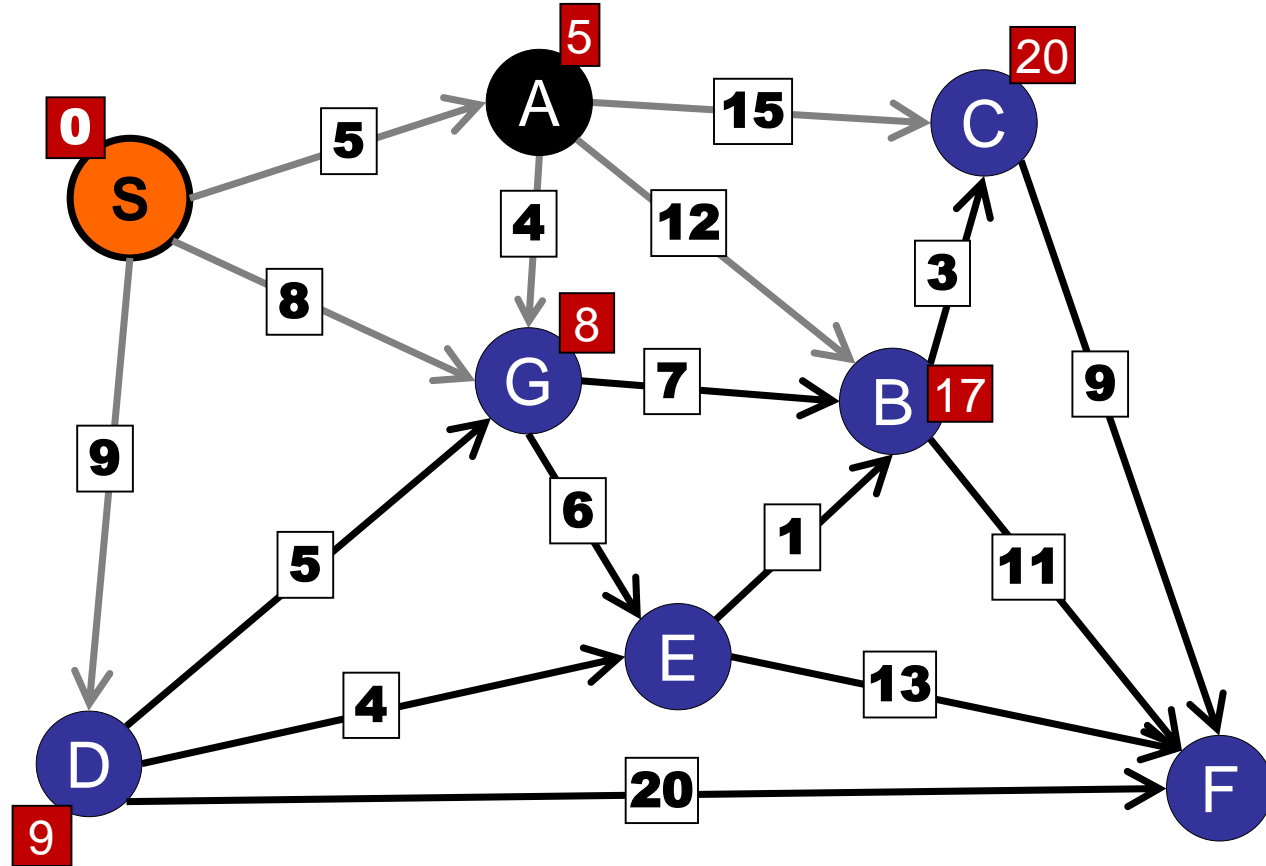
Step 2: Remove S and relax.



Vertex	Dist.
A	5
G	8
D	9

# Dijkstra's Algorithm

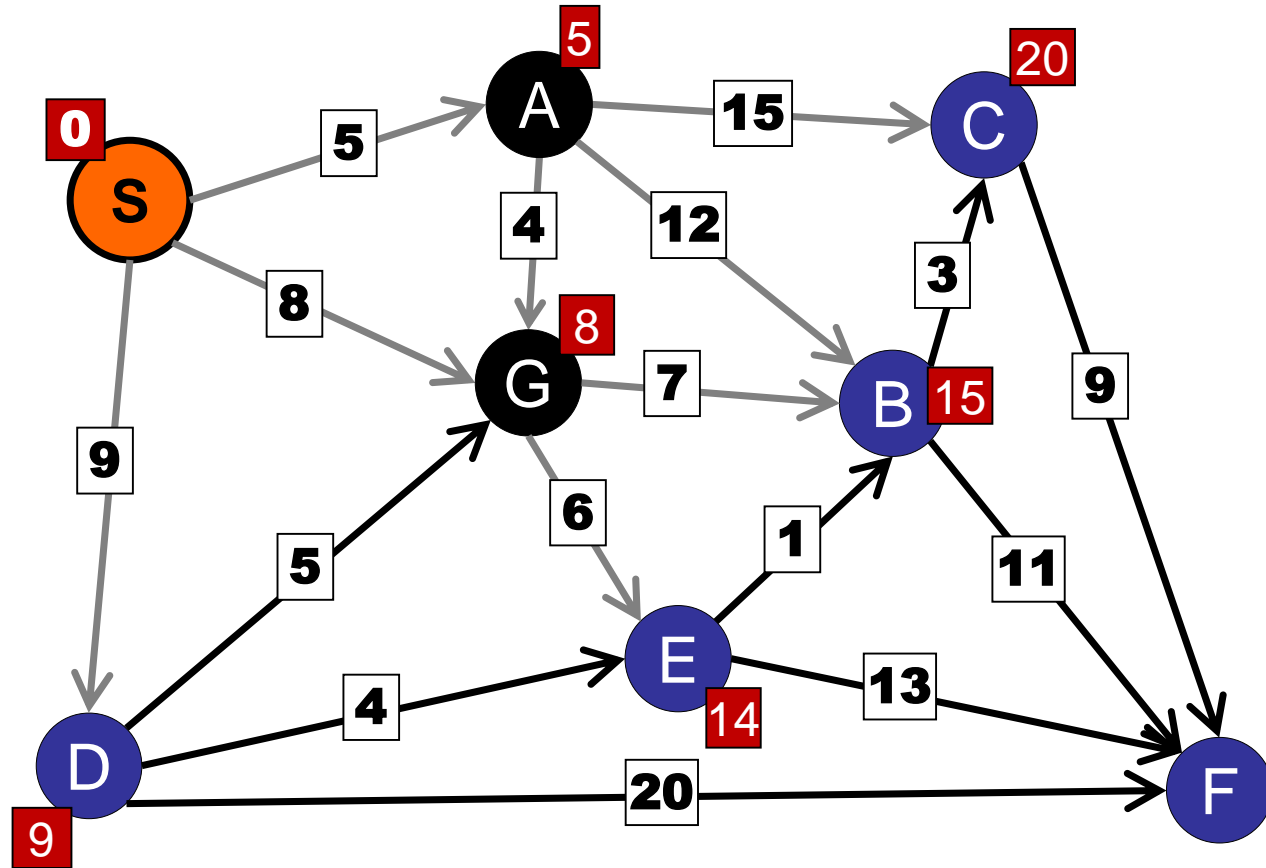
Step 3: Remove A and relax.



Vertex	Dist.
G	8
D	9
B	17
C	20

# Dijkstra's Algorithm

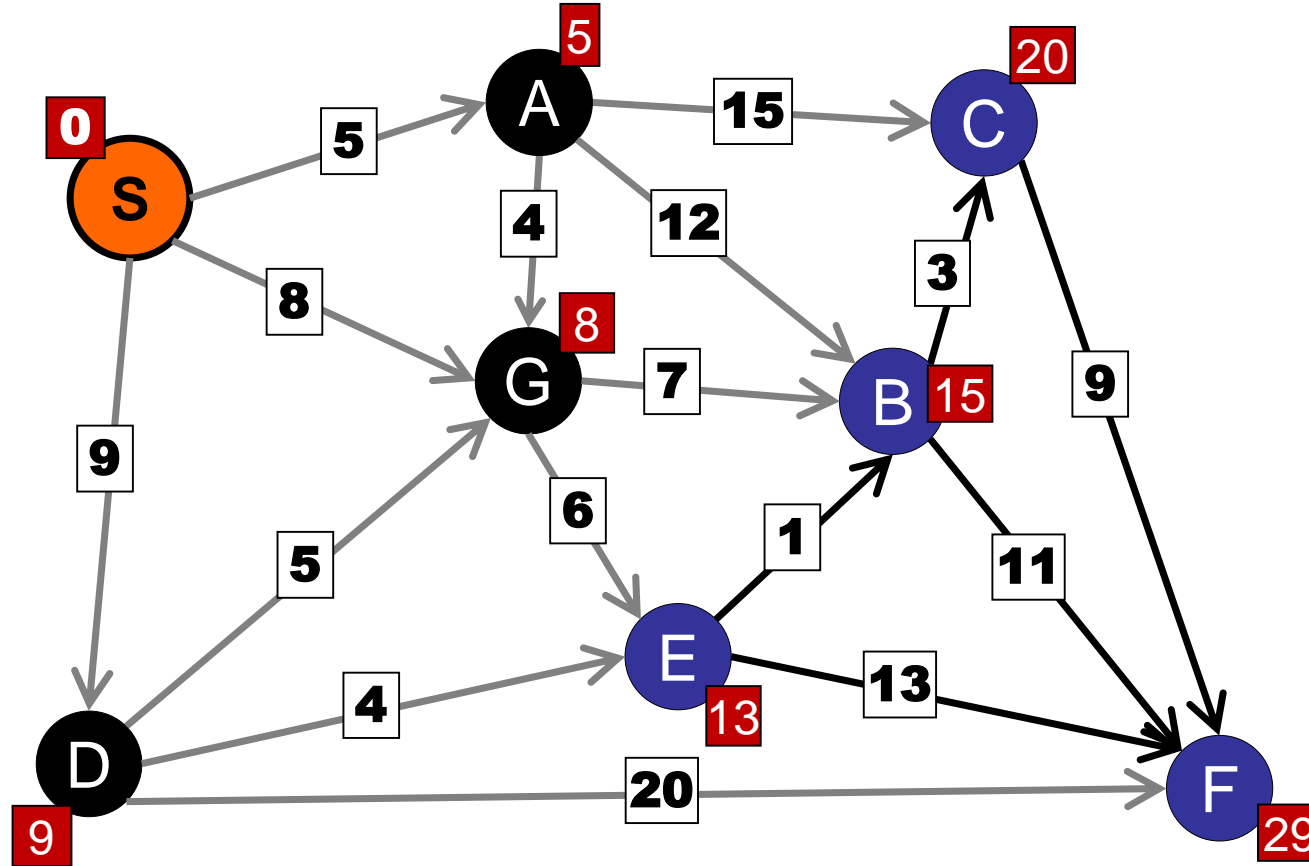
Step 4: Remove G and relax.



Vertex	Dist.
D	9
E	14
B	15
C	20

# Dijkstra's Algorithm

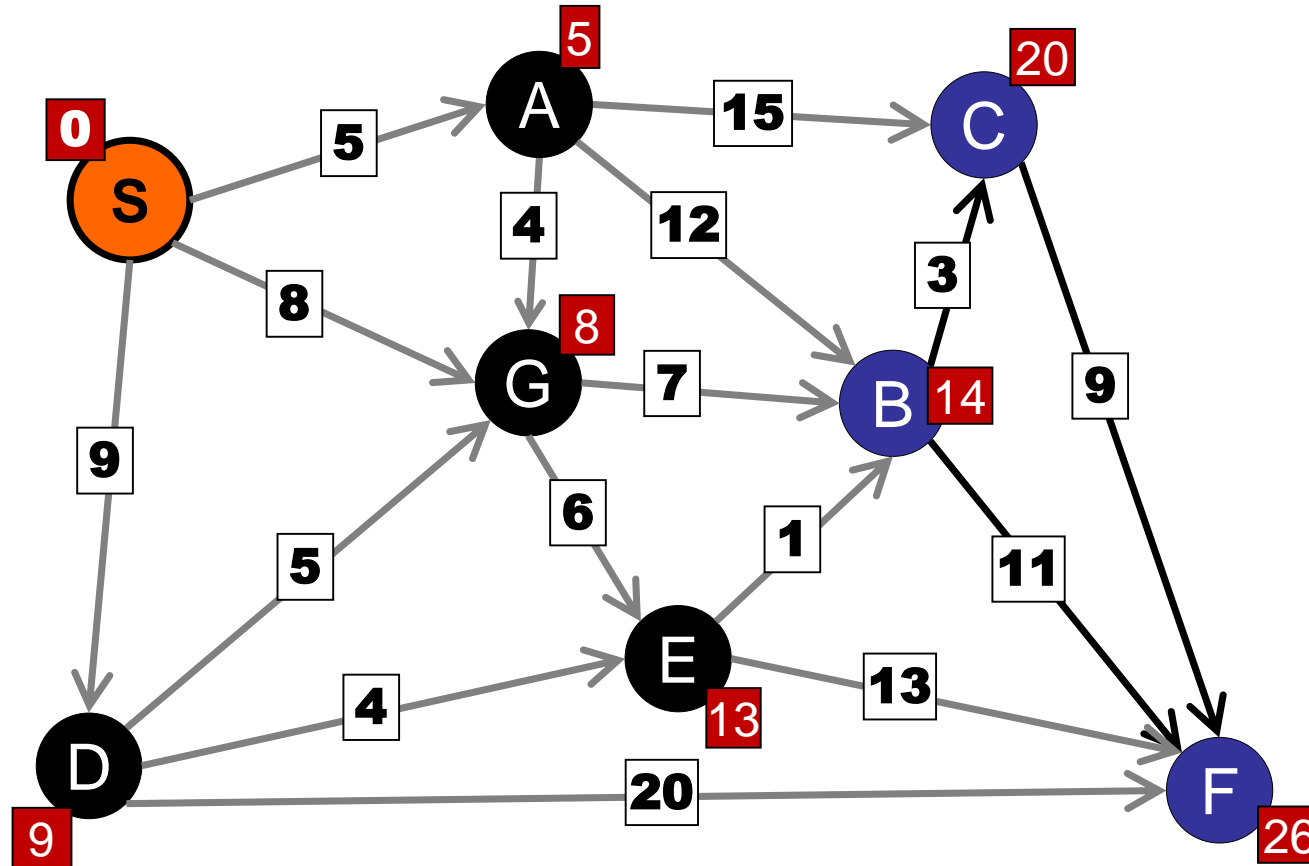
Step 5: Remove D and relax.



Vertex	Dist.
E	13
B	15
C	20
F	29

# Dijkstra's Algorithm

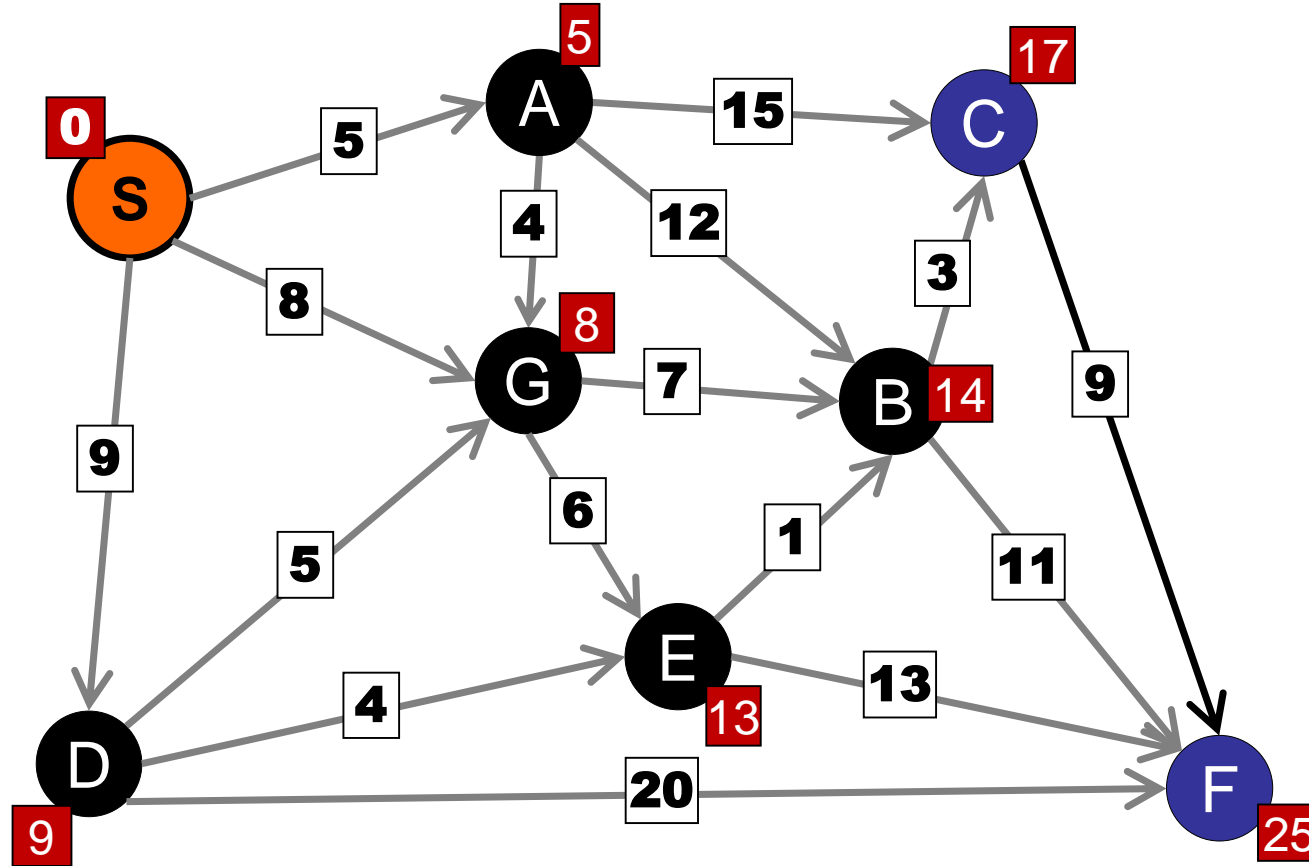
Step 5: Remove E and relax.



Vertex	Dist.
B	14
C	20
F	26

# Dijkstra's Algorithm

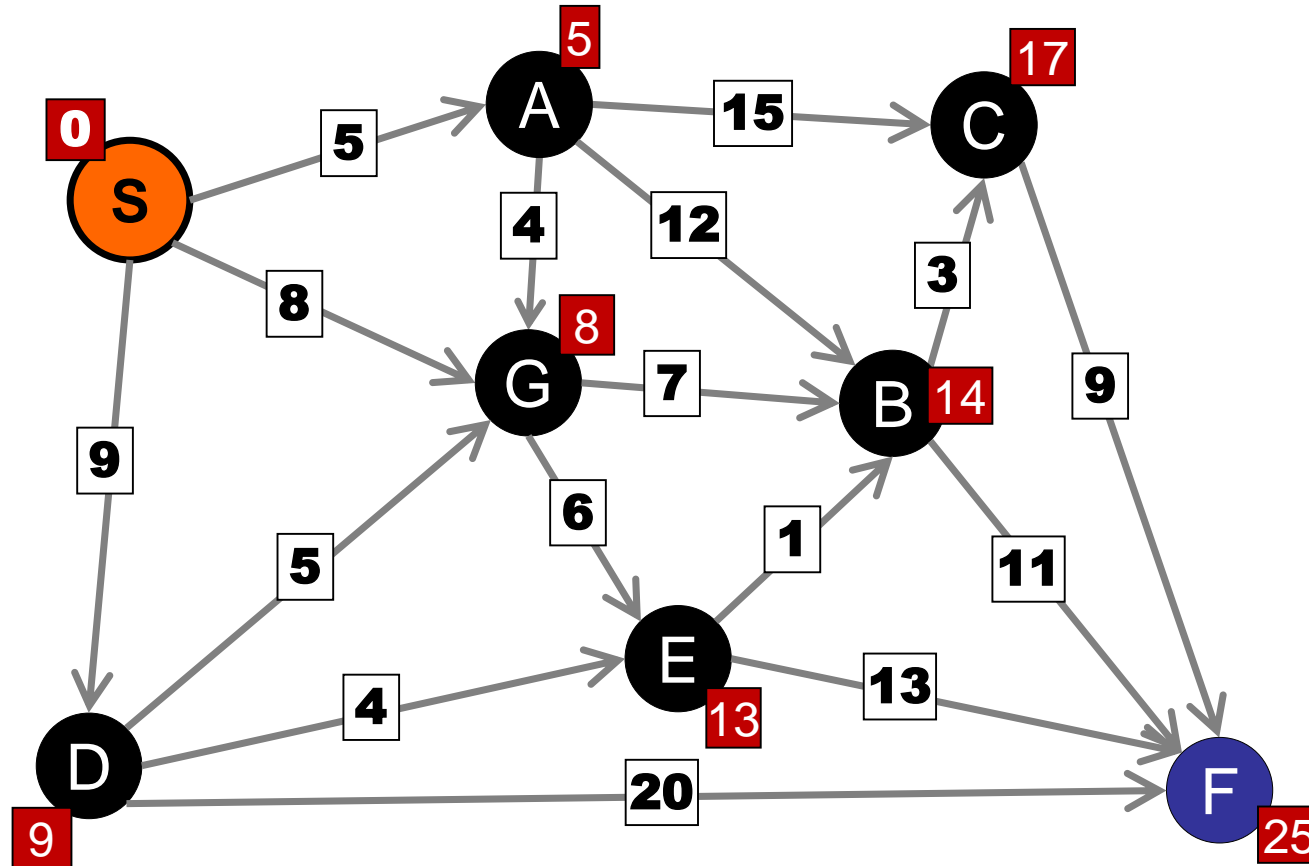
Step 5: Remove B and relax.



Vertex	Dist.
C	20
F	25

# Dijkstra's Algorithm

Step 5: Remove C and relax.



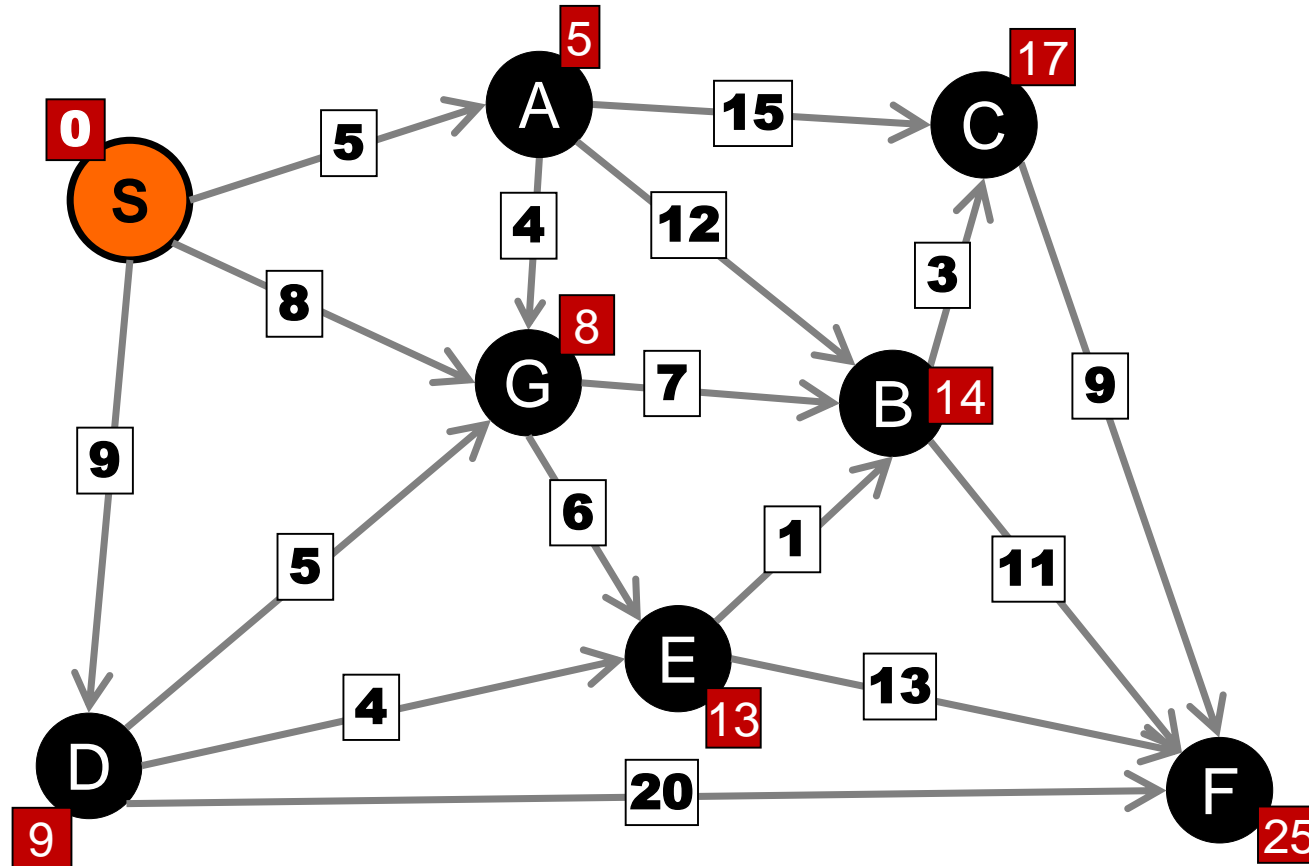
Vertex	Dist.
F	25



# Dijkstra's Algorithm

Step 5: Remove F and relax.

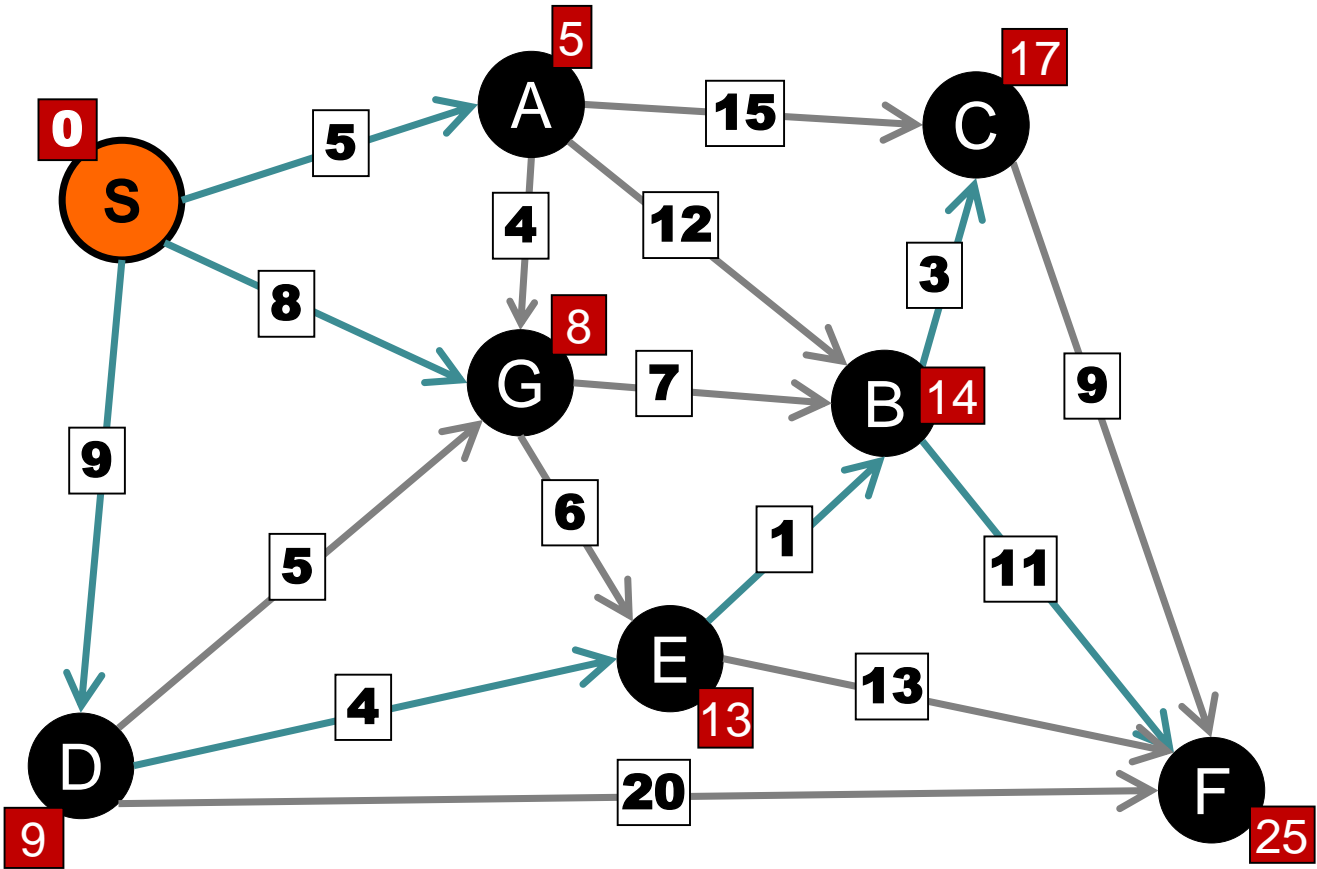
Vertex	Dist.



# Dijkstra's Algorithm

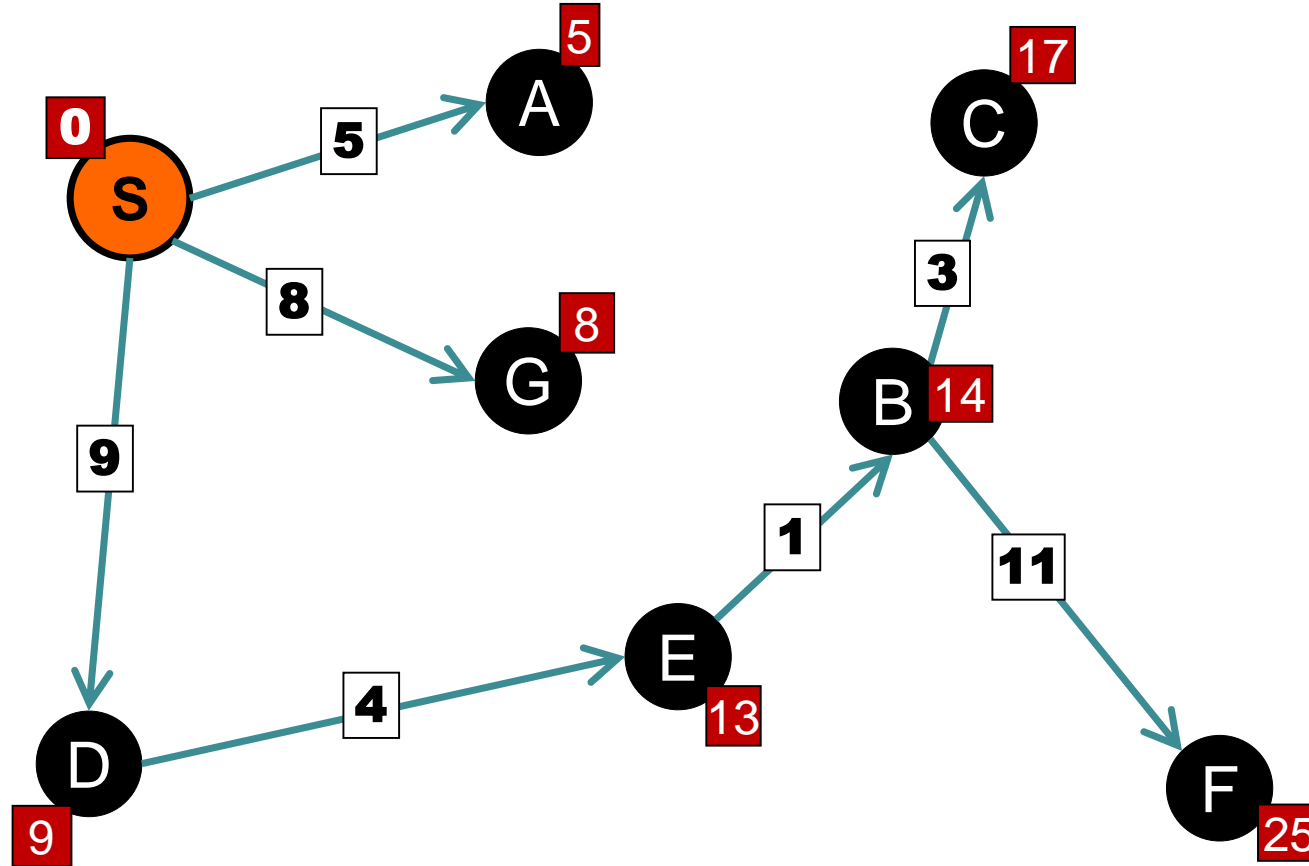
Done

Vertex	Dist.



# Dijkstra's Algorithm

## Shortest Path Tree



Vertex	Dist.

# Abstract Data Type

---

## Priority Queue

```
interface    IPriorityQueue<Key, Priority>
```

```
    void    insert(Key k, Priority p)
```

*insert k with  
priority p*

```
    Data    extractMin()
```

*remove key with  
minimum priority*

```
    void    decreaseKey(Key k, Priority p)
```

*reduce the priority of  
key k to priority p*

```
    boolean contains(Key k)
```

*does the priority  
queue contain key k?*

```
    boolean isEmpty()
```

*is the priority queue  
empty?*

### Notes:

Assume data items are unique.

```
public Dijkstra{
    private Graph G;
    private IPriorityQueue pq = new PriQueue();
    private double[] distTo;

    searchPath(int start) {
        pq.insert(start, 0.0);
        distTo = new double[G.size()];
        Arrays.fill(distTo, INFTY);
        distTo[start] = 0;
        while (!pq.isEmpty()) {
            int w = pq.deleteMin();
            for (Edge e : G[w].nbrList)
                relax(e);
        }
    }
}
```

...

# Dijkstra's Algorithm

---

```
relax(Edge e) {  
    int v = e.from();  
    int w = e.to();  
    double weight = e.weight();  
    if (distTo[w] > distTo[v] + weight) {  
        distTo[w] = distTo[v] + weight;  
        parent[w] = v;  
        if (pq.contains(w))  
            pq.decreaseKey(w, distTo[w]);  
        else  
            pq.insert(w, distTo[w]);  
    }  
}
```

# Abstract Data Type

---

## Priority Queue

```
interface    IPriorityQueue<Key, Priority>
```

```
    void    insert(Key k, Priority p)
```

*insert k with  
priority p*

```
    Data    extractMin()
```

*remove key with  
minimum priority*

```
    void    decreaseKey(Key k, Priority p)
```

*reduce the priority of  
key k to priority p*

```
    boolean contains(Key k)
```

*does the priority  
queue contain key k?*

```
    boolean isEmpty()
```

*is the priority queue  
empty?*

### Notes:


Assume data items are unique.

# Priority Queue

---

## AVL Tree

- Indexed by: priority
- Existing operations:
  - deleteMin()
  - insert(key, priority)






# Priority Queue

---

## AVL Tree

- Other operations:
  - contains(key)
  - decreaseKey(key, priority)
- How to find a vertex?



# Priority Queue

---


## AVL Tree

### – Other operations:

- contains(key)
- decreaseKey(key, priority)

### – Hash Table:

- Map keys to location in AVL tree.
- Update hash table whenever the binary tree changes.



G: 9
B: 15
H: 13
C: 22
I: 28
A: 5
D: 6
E: 1
F: 7

# Dijkstra's Algorithm

---

## Priority Queue by AVL tree:

- **insert(key, priority):  $O(\log n)$** 
  - Insert (priority, key) in AVL tree indexed by priority
  - Insert (key, priority) in hash table
- **deleteMin():  $O(\log n)$** 
  - Find node with the minimum priority and delete it from AVL tree
- **decreaseKey(key, priority):  $O(\log n)$** 
  - Find current priority (curPri) of key in hash table, remove (curPri, key) from AVL tree, insert (priority, key) into AVL tree, update hash table record for key
- **contains(key):  $O(1)$** 
  - Search in the hash table for key

# Dijkstra's Algorithm

---

## Priority Queue by AVL tree:

- **insert(key, priority):  $O(\log n)$** 
  - Insert (priority, key) in AVL tree indexed by priority
  - Insert (key, priority) in hash table
- **deleteMin():  $O(\log n)$** 
  - Find node with the minimum priority and delete it
- **decreaseKey(key, priority):  $O(\log n)$** 
  - Find current priority (curPri) of key in hash table, remove (curPri, key) from AVL tree, insert (priority, key) into AVL tree, update hash table record for key
- **contains(key):  $O(1)$** 
  - Search in the hash table for key

What if there are multiple keys with same priority? Possible approaches:

- Put all keys with same priority in the same node in AVL tree
- Have distinct nodes in AVL tree for different keys, but in hash table, include pointer to the particular node for each key.

What is the running time of Dijkstra's Algorithm, using an AVL tree Priority Queue?

1.  $O(V + E)$
- ✓ 2.  $O(E \log V)$
3.  $O(V \log E)$
4.  $O(V^2)$
5.  $O(VE)$
6. None of the above



```
public Dijkstra{
    private Graph G;
    private MinPriQueue pq = new MinPriQueue();
    private double[] distTo;

    searchPath(int start) {
        pq.insert(start, 0.0);
        distTo = new double[G.size()];
        Arrays.fill(distTo, INFTY);
        distTo[start] = 0;
        while (!pq.isEmpty()) {
            int w = pq.deleteMin();
            for (Edge e : G[w].nbrList)
                relax(e);
        }
    }
}
```

How many times?

How many times?

# Dijkstra's Algorithm

---

```
relax(Edge e) {  
    int v = e.from();  
    int w = e.to();  
    double weight = e.weight();  
    if (distTo[w] > distTo[v] + weight) {  
        distTo[w] = distTo[v] + weight;  
        parent[w] = v;  
        if (pq.contains(w))  
            pq.decreaseKey(w, distTo[w]);  
        else  
            pq.insert(w, distTo[w]);  
    }  
}
```

# Dijkstra's Algorithm

---

## Analysis:

- insert / deleteMin:  $|V|$  times each
  - Each node is added to the priority queue **once**.
- relax / decreaseKey:  $|E|$  times
  - Each edge is relaxed once.
- Priority queue operations:  $O(\log V)$
- Total:  $O((V+E)\log V) = O(E \log V)$



# Dijkstra's Algorithm


---

Why does it work?

# Dijkstra's Algorithm

---

Every edge crossing the boundary has been relaxed.




**finished vertices:**  
**distance is accurate.**  
**Initially: just the source.**

# Dijkstra's Algorithm

---

Every edge crossing the boundary has been relaxed.

**finished vertices:  
distance is accurate.**




**fringe vertices:  
neighbor of a  
finished vertex.**

**fringe vertices: in priority queue  
neighbor of a finished vertex.**

# Dijkstra's Algorithm

Every edge crossing the boundary has been relaxed.

finished vertices:  
distance is accurate.



fringe vertices:  
neighbor of a  
finished vertex.

other vertices:  
no known  
estimate

fringe vertices: in priority queue  
neighbor of a finished vertex.

# Dijkstra's Algorithm

---

Proof by induction:

- Every “finished” vertex has correct estimate.
- Initially: only “finished” vertex is start.

# Dijkstra's Algorithm

---


Proof by induction:

- Every “finished” vertex has correct estimate.
- Initially: only “finished” vertex is start.
- Inductive step:
  - Remove vertex from priority queue.
  - Relax its edges.
  - Add it to finished.
  - **Claim: it has a correct estimate.**

# Dijkstra's Algorithm

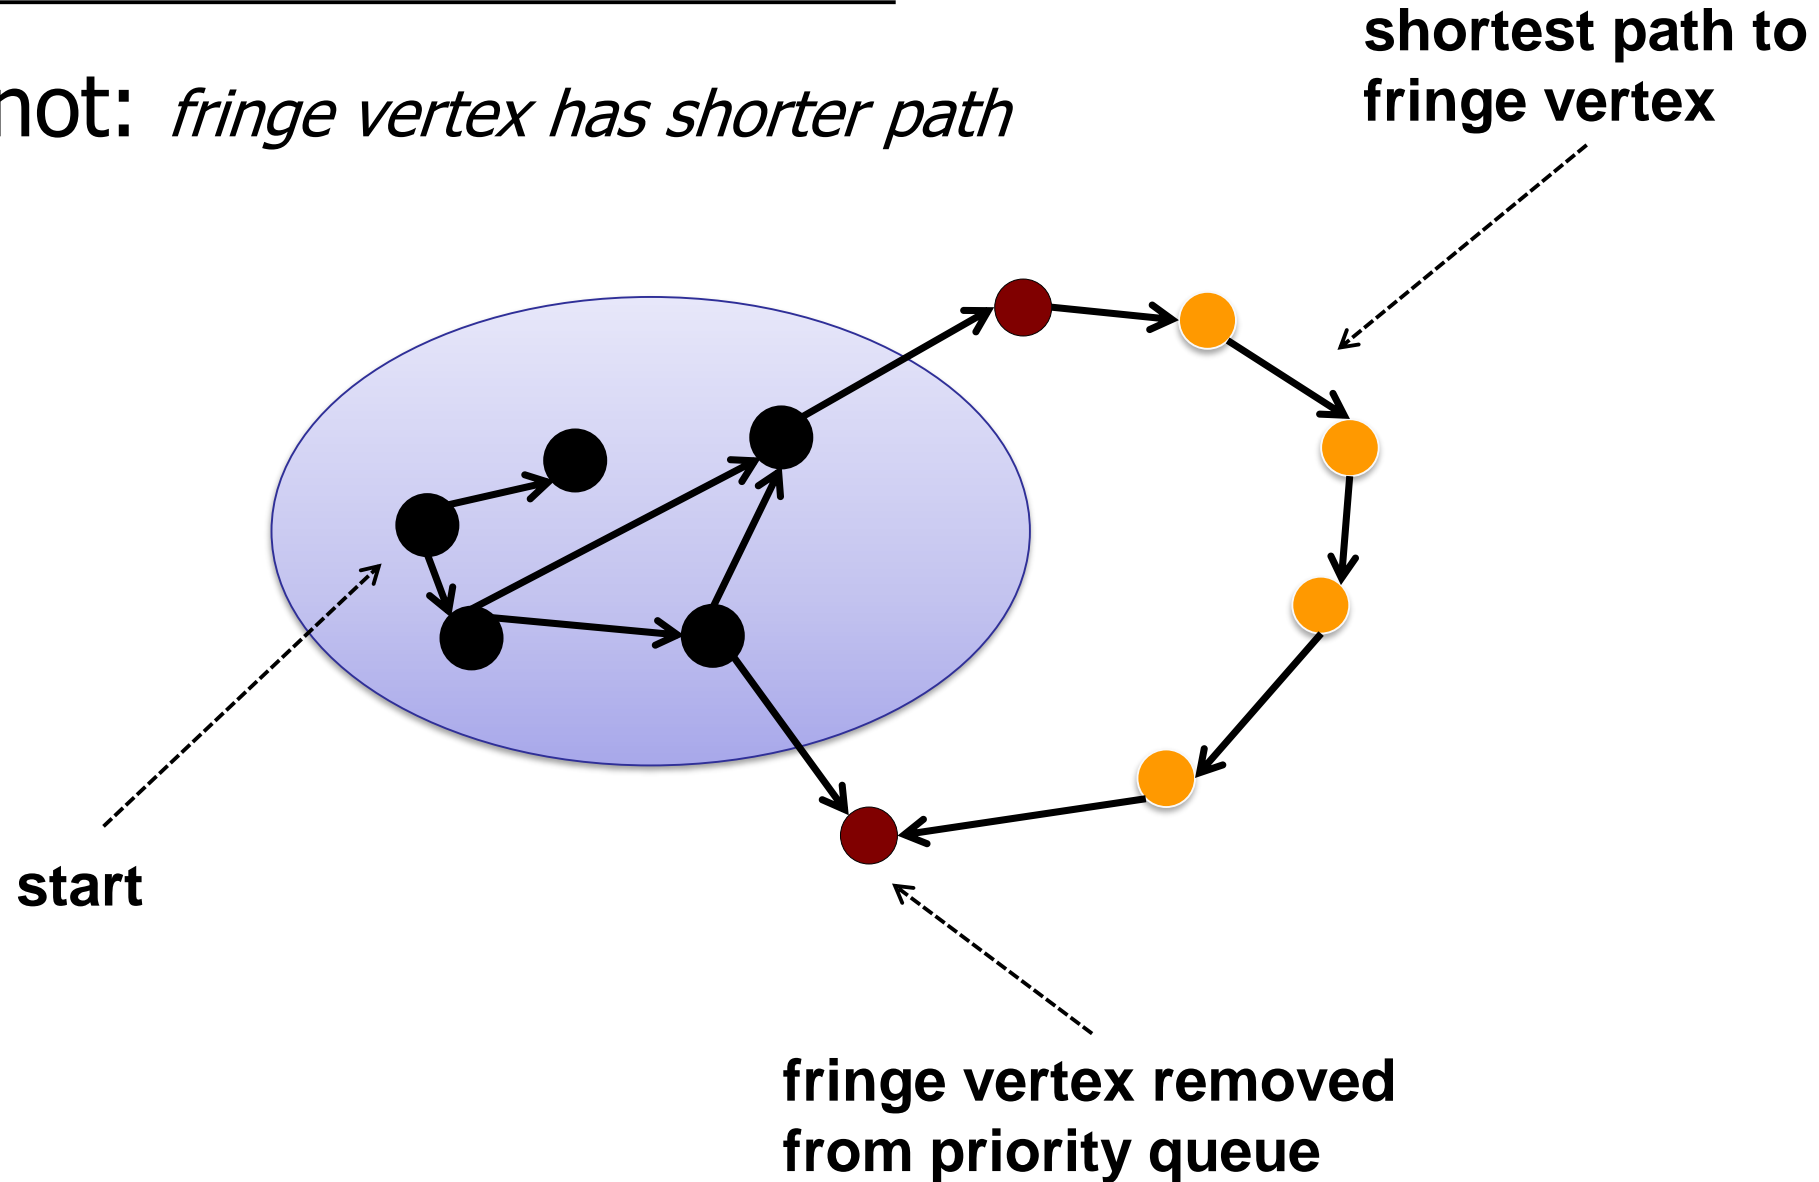
---

Assume not: *fringe vertex is removed but not done*



# Dijkstra's Algorithm

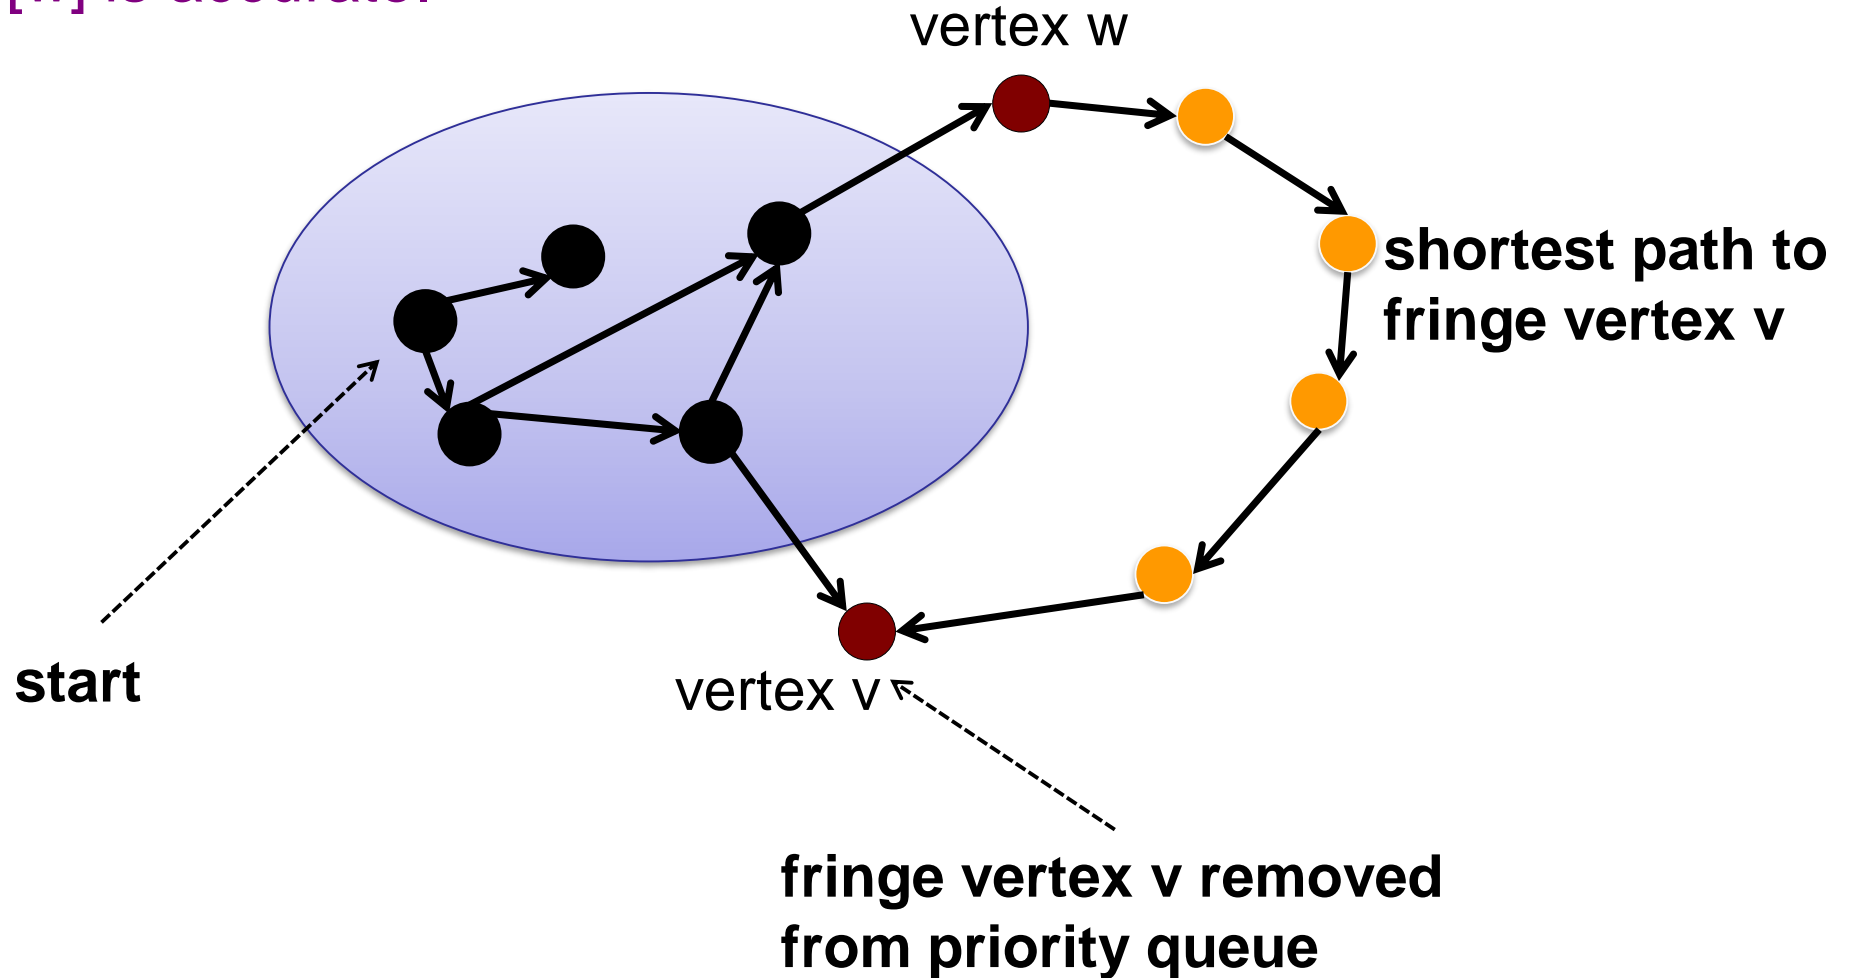
Assume not: *fringe vertex has shorter path*





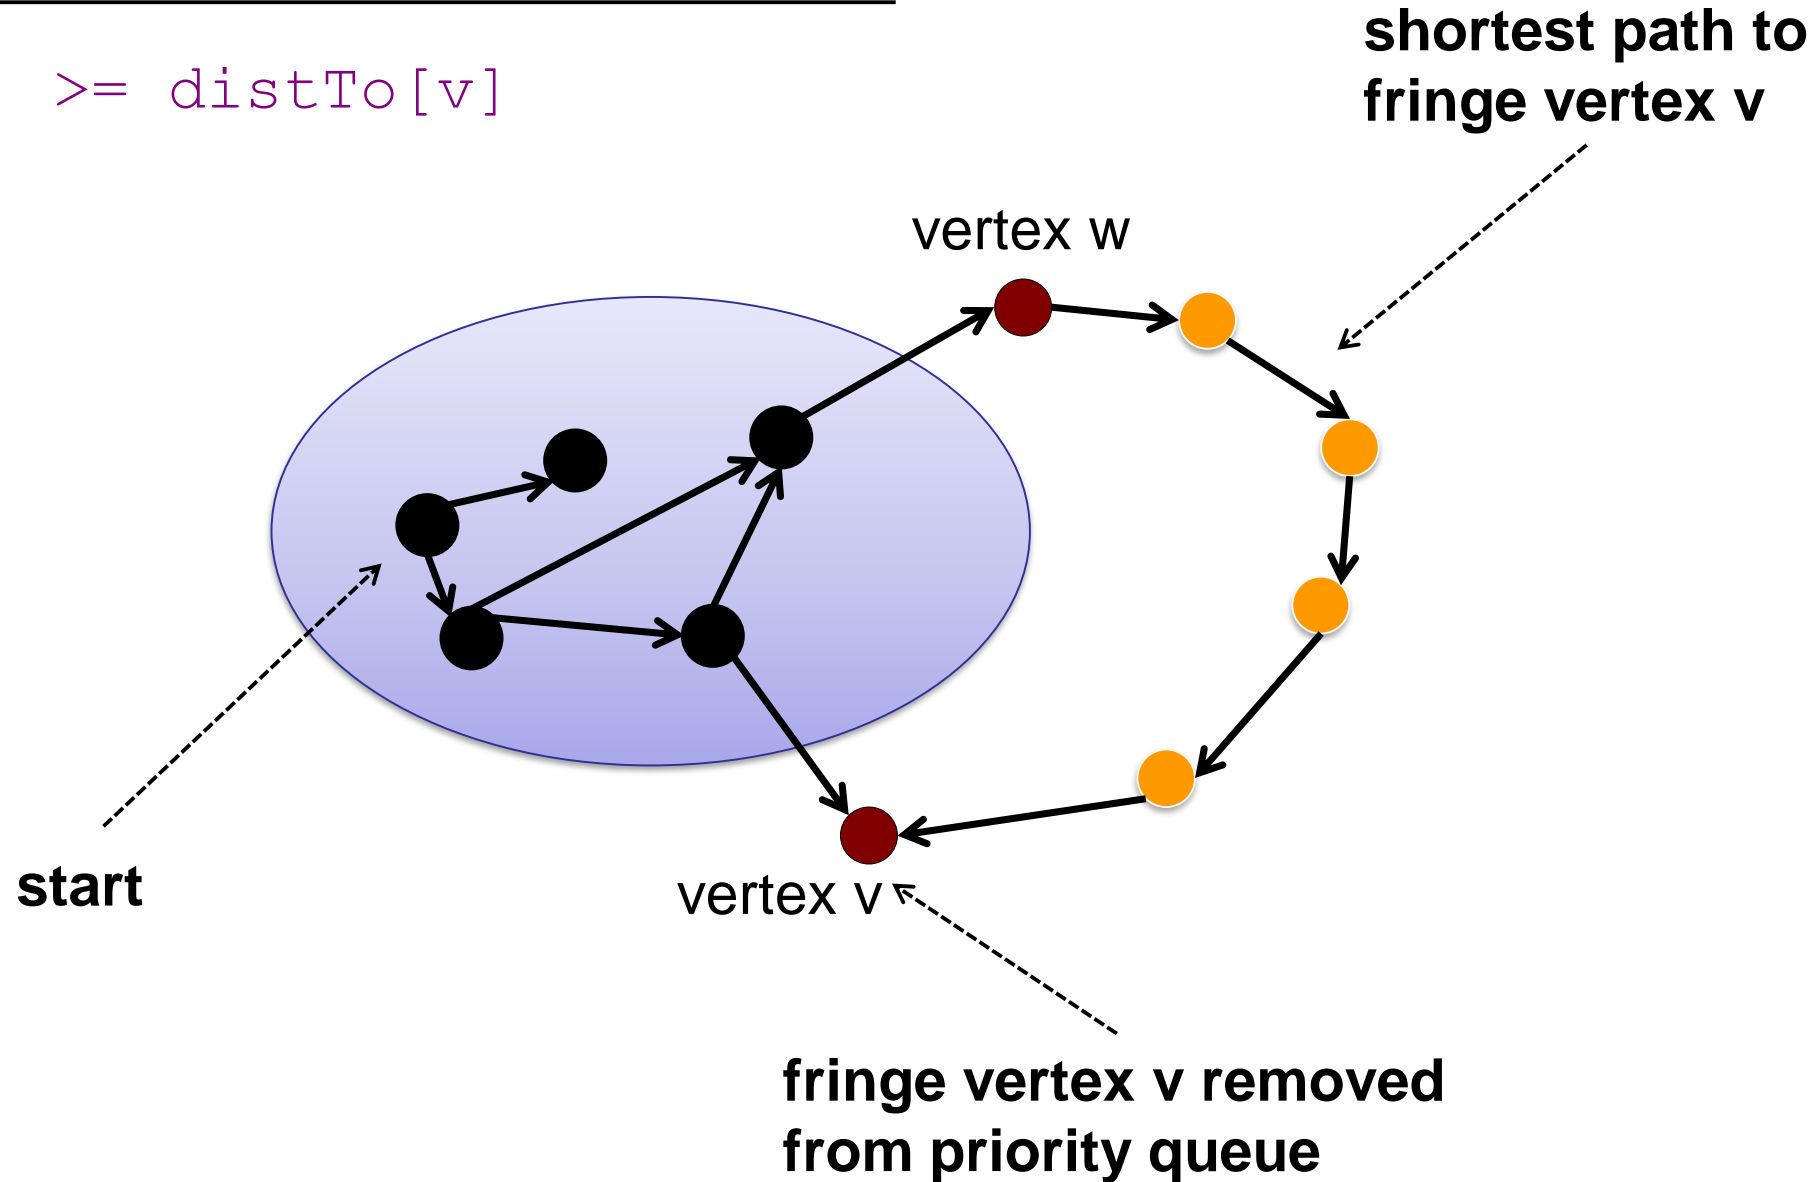
# Dijkstra's Algorithm

If  $P$  is shortest path to  $v$ , then prefix of  $P$  is shortest path to  $w$ .  
Then  $\text{distTo}[w]$  is accurate.



# Dijkstra's Algorithm

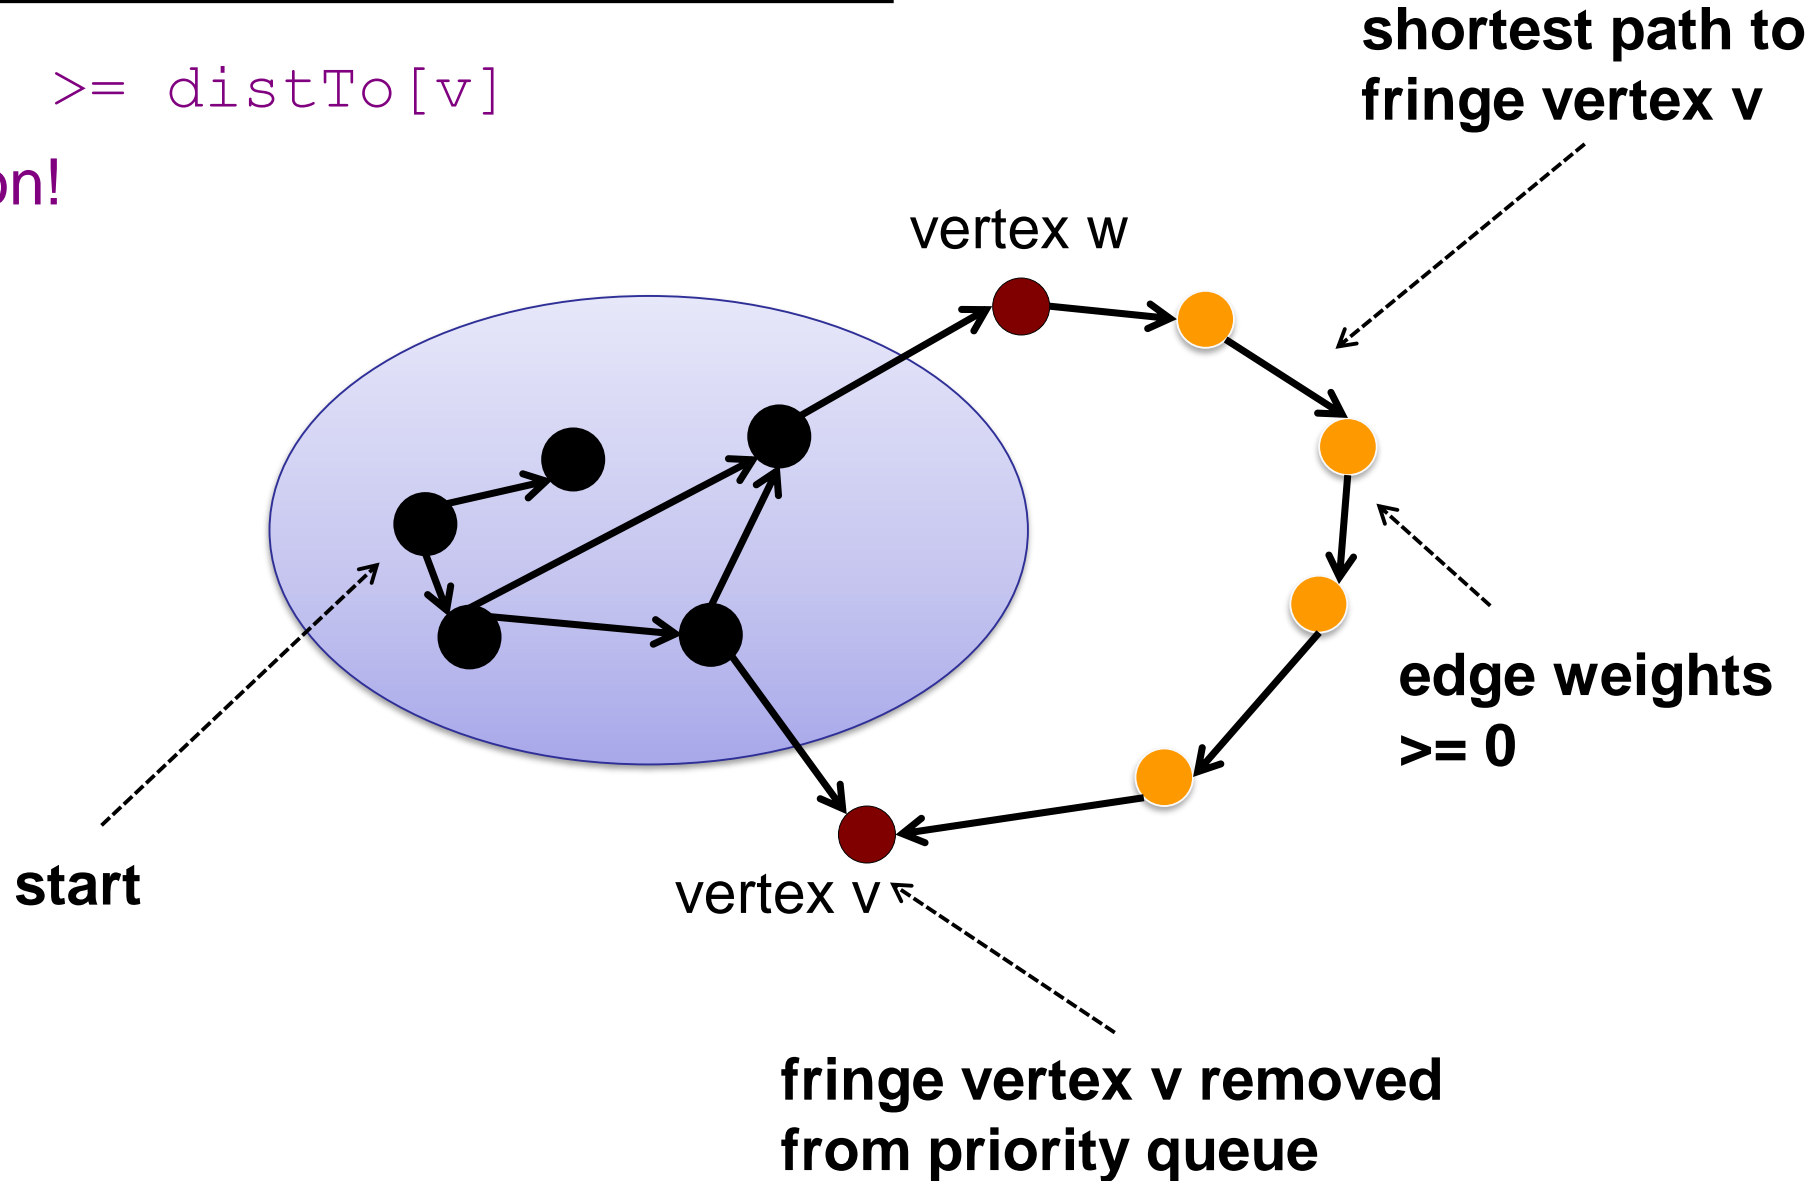
$\text{distTo}[w] \geq \text{distTo}[v]$



# Dijkstra's Algorithm

$\text{distTo}[w] \geq \text{distTo}[v]$

Contradiction!



# Dijkstra's Algorithm

---

Proof by induction:

- Every “finished” vertex has correct estimate.
- Initially: only “finished” vertex is start.
- Inductive step:
  - Remove vertex from priority queue.
  - Relax its edges.
  - Add it to finished.
  - **Claim: it has a correct estimate.**

# Dijkstra's Algorithm

---

```
relax(Edge e) {  
    int v = e.from();  
    int w = e.to();  
    double weight = e.weight();  
    if (distTo[w] > distTo[v] + weight) {  
        distTo[w] = distTo[v] + weight;  
        parent[w] = v;  
        if (pq.contains(w))  
            pq.decreaseKey(w, distTo[w]);  
        else  
            pq.insert(w, distTo[w]);  
    }  
}
```

**Extending a path does not make it shorter!**

# Dijkstra's Algorithm

---

## Analysis:

- insert / deleteMin:  $|V|$  times each
  - Each node is added to the priority queue **once**.
- decreaseKey:  $|E|$  times
  - Each edge is relaxed once.
- Priority queue operations:  $O(\log V)$
- Total:  $O((V+E)\log V) = O(E \log V)$

## Source-to-Destination Dijkstra

Can we stop as soon as we dequeue the destination?

- ✓ 1. Yes.
- 2. Only if the graph is sparse.
- 3. No.

# Dijkstra's Algorithm

---

## Source-to-Destination:

- What if you stop the first time you dequeue the destination?
- Recall:
  - a vertex is “finished” when it is dequeued
  - the estimate is *never* changed again
  - if the destination is finished, then stop



# Dijkstra Summary

---

Basic idea:

- Maintain distance estimates.
- Repeat:
  - Find unfinished vertex with smallest estimate.
  - Relax all outgoing edges.
  - Mark vertex finished.
- $O(E \log V)$  time (with AVL tree).

# Dijkstra's Performance

---

PQ Implementation	insert	deleteMin	decreaseKey	Total
Array	1	V	1	<b><math>O(V^2)</math></b>
AVL Tree	$\log V$	$\log V$	$\log V$	<b><math>O(E \log V)</math></b>
d-way Heap	$d \log_d V$	$d \log_d V$	$\log_d V$	<b><math>O(E \log_{E/V} V)</math></b>
Fibonacci Heap	1	$\log V$	1	<b><math>O(E + V \log V)</math></b>

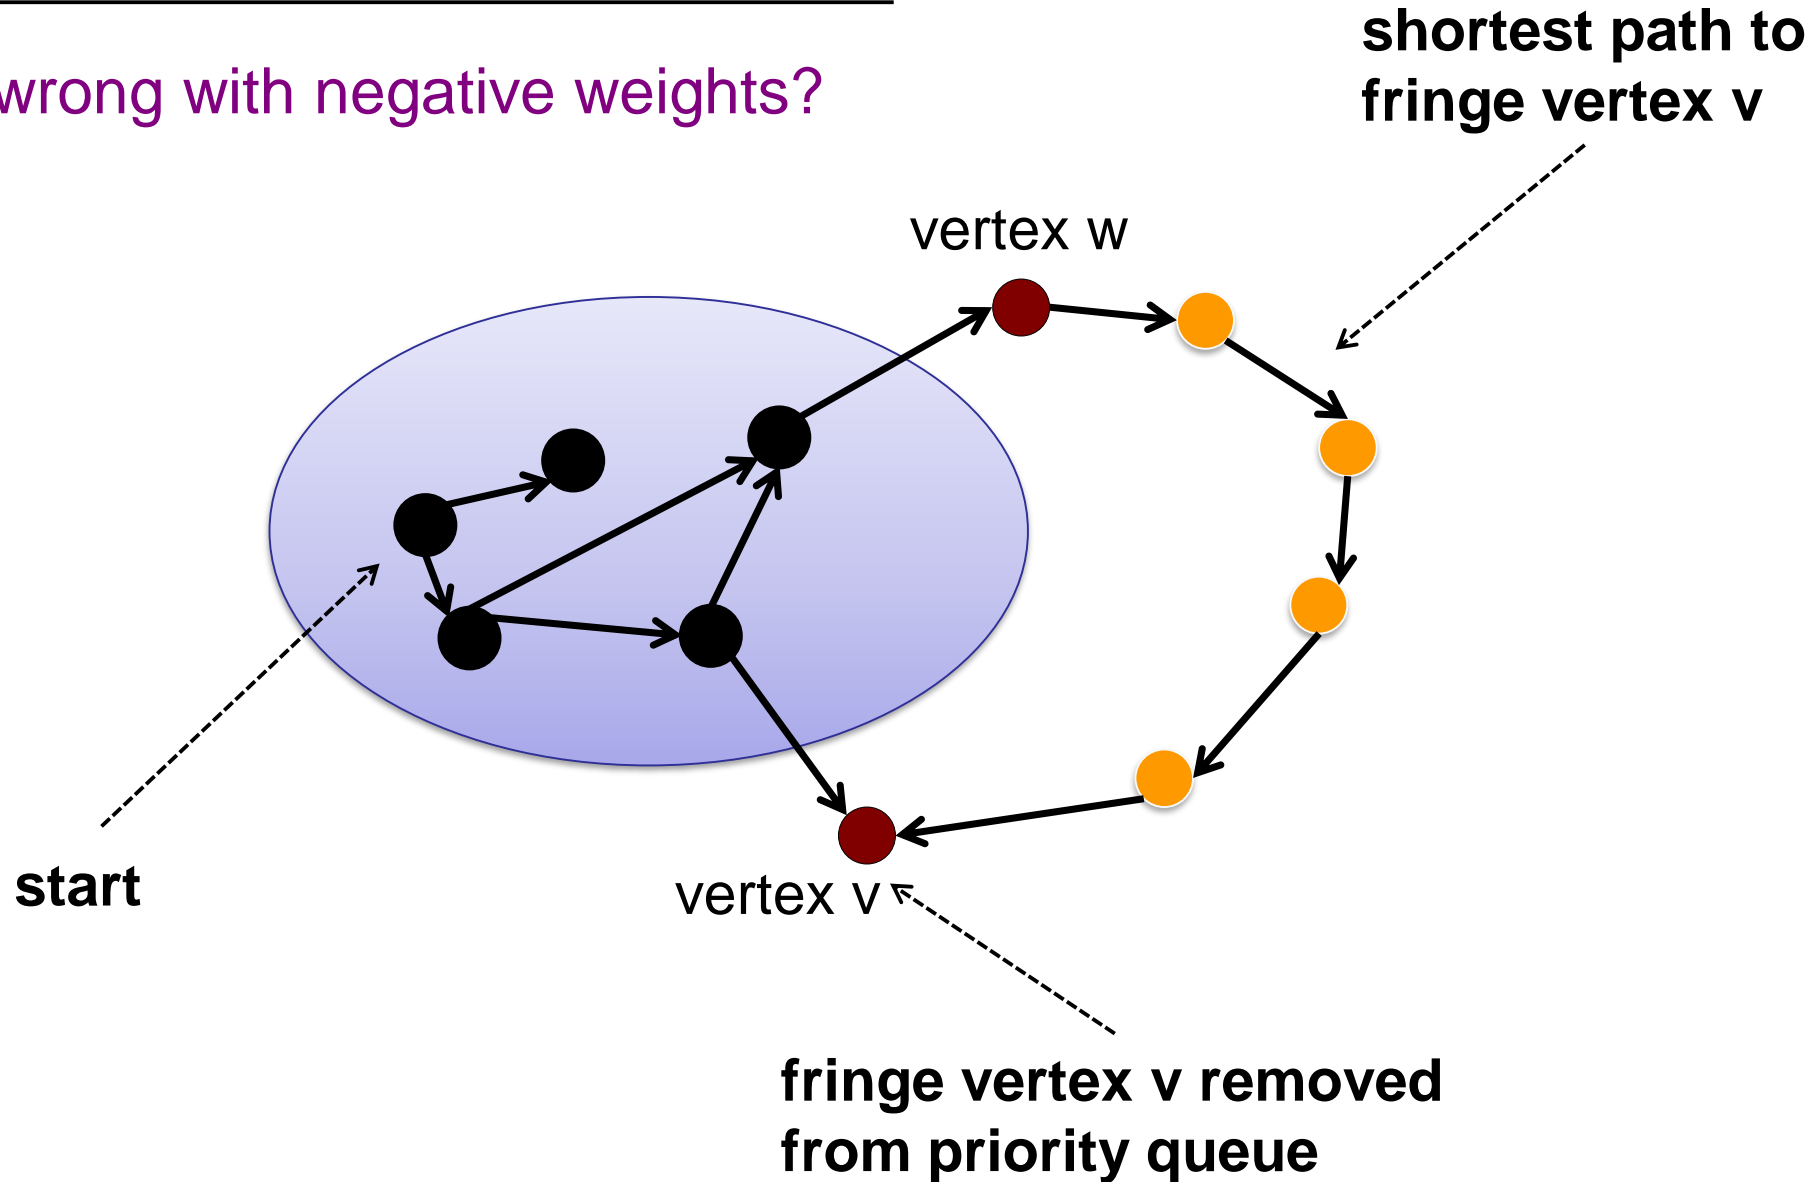
# Dijkstra Summary

---

Edges with negative weights?

# Dijkstra's Algorithm


What goes wrong with negative weights?



# Dijkstra's Algorithm

---


Edges with negative weights?



# Dijkstra's Algorithm

---

Edges with negative weights?




Step 1: Remove A.  
Relax A.  
Mark A done.

# Dijkstra's Algorithm

---

Edges with negative weights?



Step 1: Remove A.  
Relax A.  
Mark A done.


...

Step 4: Remove B.  
Relax B.  
Mark B done.

Oops: We need to  
update A.

# Dijkstra's Algorithm

What goes wrong with negative weights?






# Dijkstra's Algorithm

---

Can we reweight?

e.g.:  $\text{weight} += 10$




# Can we reweight the graph?

1. Yes.
2. Only if there are no negative weight cycles.
- ✓ 3. No.

# Dijkstra's Algorithm

---

Can we reweight?




Path S-B-A: 1

Path S-A: 4

# Dijkstra's Algorithm

---

Can we reweight?



Path S-B-A: 21

Path S-A: 14

# Dijkstra Summary

---

## Basic idea:

- Maintain distance estimates.
- Repeat:
  - Find unfinished vertex with smallest estimate.
  - Relax all outgoing edges.
  - Mark vertex finished.
- $O(E \log V)$  time (with AVL tree Priority Queue).
- No negative weight edges!

# Dijkstra Comparison

---

Same algorithm:

- Maintain a set of explored vertices.
  - Add vertices to the explored set by following edges that go from a vertex in the explored set to a vertex outside the explored set.
- 
- **BFS**: Take edge from vertex that was discovered **least** recently.
  - **DFS**: Take edge from vertex that was discovered **most** recently.
  - **Dijkstra's**: Take edge from vertex that is **closest** to source.

# Dijkstra Comparison

---

Same algorithm:

- Maintain a set of explored vertices.
  - Add vertices to the explored set by following edges that go from a vertex in the explored set to a vertex outside the explored set.
- 
- BFS: Use queue.
  - DFS: Use stack.
  - Dijkstra's: Use priority queue.

# What about for Negative Weights?

---

- Only in October '22, an algorithm solving SSSP with negative weight edges with  $\tilde{O}(E)$  running time proposed!

## Negative-Weight Single-Source Shortest Paths in Near-linear Time

Aaron Bernstein\*

Danupon Nanongkai<sup>†</sup>

Christian Wulff-Nilsen<sup>‡</sup>

### Abstract

We present a randomized algorithm that computes single-source shortest paths (SSSP) in  $O(m \log^8(n) \log W)$  time when edge weights are integral and can be negative.<sup>1</sup> This essentially resolves the classic negative-weight SSSP problem. The previous bounds are  $\tilde{O}((m + n^{1.5}) \log W)$  [BLNPSSW FOCS'20] and  $m^{4/3+o(1)} \log W$  [AMV FOCS'20]. Near-linear time algorithms were known previously only for the special case of planar directed graphs [Fakcharoenphol and Rao FOCS'01].

In contrast to all recent developments that rely on sophisticated continuous optimization methods and dynamic algorithms, our algorithm is simple: it requires only a simple graph decomposition and elementary combinatorial tools. In fact, ours is the first combinatorial algorithm for negative-weight SSSP to break through the classic  $\tilde{O}(m\sqrt{n} \log W)$  bound from over three decades ago [Gabow and Tarjan SICOMP'89].